


# Apache Spark

Adapted from the slides by Dr. Zareen Alamgir

*Sources for Slides are Spark 2.0.0 Documentation, Learning Spark(BOOK), Jimmy Lin notes, DataBricks (Holden Karau), Spark Tutorial by tutorialpoint*



# Spark vs. Hadoop



Google Trends

# Why Spark ?

Hadoop is extensively used to analyze data sets

- It provides computing solution that is scalable, flexible, fault-tolerant and cost effective.

Why Spark ?

- Fast Computations (in-memory cluster computing)
- Iterative Algorithms
- Interactive Queries
- Stream Processing
- Graph Algorithms
- Reduce management burden of maintaining separate tools.

# Why Spark ?

Hadoop is great, but it's  
really way too low  
level!

*(circa 2007)*

What's the  
solution?

Design a  
higher-level  
language

# What's Spark

Apache Spark is a fast cluster computing platform designed for fast computation.

Built on top of Hadoop MapReduce and extends it to efficiently use more types of computations

- Interactive Queries
- Stream Processing

**Up to 100 times faster in memory, and 10 times faster when running on disk.**

# Difficulty of Programming in MR

## Word Count implementations

- Hadoop MR – 61 lines in Java
- Spark – 1 line in interactive shell

```
sc.textFile('...').flatMap(lambda x: x.split())  
    .map(lambda x: (x, 1)).reduceByKey(lambda x, y: x+y)  
    .saveAsTextFile('...')
```

VS

```
import java.io.IOException;  
import java.util.StringTokenizer;  
  
import org.apache.hadoop.conf.Configuration;  
import org.apache.hadoop.fs.Path;  
import org.apache.hadoop.io.IntWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Job;  
import org.apache.hadoop.mapreduce.Mapper;  
import org.apache.hadoop.mapreduce.Reducer;  
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;  
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;  
  
public class WordCount {  
  
    public static class TokenizerMapper  
        extends Mapper<Object, Text, Text, IntWritable>{  
  
        private final static IntWritable one = new IntWritable(1);  
        private Text word = new Text();  
  
        public void map(Object key, Text value, Context context  
            ) throws IOException, InterruptedException {  
            StringTokenizer itr = new StringTokenizer(value.toString());  
            while (itr.hasMoreTokens()) {  
                word.set(itr.nextToken());  
                context.write(word, one);  
            }  
        }  
    }  
  
    public static class IntSumReducer  
        extends Reducer<Text, IntWritable, Text, IntWritable> {  
        private IntWritable result = new IntWritable();  
  
        public void reduce(Text key, Iterable<IntWritable> values,  
            Context context  
            ) throws IOException, InterruptedException {  
  
            int sum = 0;  
            for (IntWritable val : values) {  
                sum += val.get();  
            }  
            result.set(sum);  
            context.write(key, result);  
        }  
    }  
  
    public static void main(String[] args) throws Exception {  
        Configuration conf = new Configuration();  
        Job job = Job.getInstance(conf, "word count");  
        job.setJarByClass(WordCount.class);  
        job.setMapperClass(TokenizerMapper.class);  
        job.setCombinerClass(IntSumReducer.class);  
        job.setReducerClass(IntSumReducer.class);  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
        FileInputFormat.addInputPath(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
        System.exit(job.waitForCompletion(true) ? 0 : 1);  
    }  
}
```

# Limitations of Map Reduce



Map Reduce is  
designed for  
Batch Processing

- Cannot handle interactive queries
- Cannot handle iterative tasks
- Cannot handle stream processing



# Spark offers you

```
graph TD; A[Spark offers you] --> B[Lazy Computations]; A --> C[In-memory data caching]; A --> D[Efficient pipelining]; B --> E[Optimize the job before executing]; C --> F[Scan HDD only once, then scan your RAM]; D --> G[Avoids the data hitting the HDD by all means];
```

Lazy  
Computations

Optimize the  
job before  
executing

In-memory data  
caching

Scan HDD only  
once, then scan  
your RAM

Efficient  
pipelining

Avoids the data  
hitting the HDD  
by all means



# Spark & Hadoop

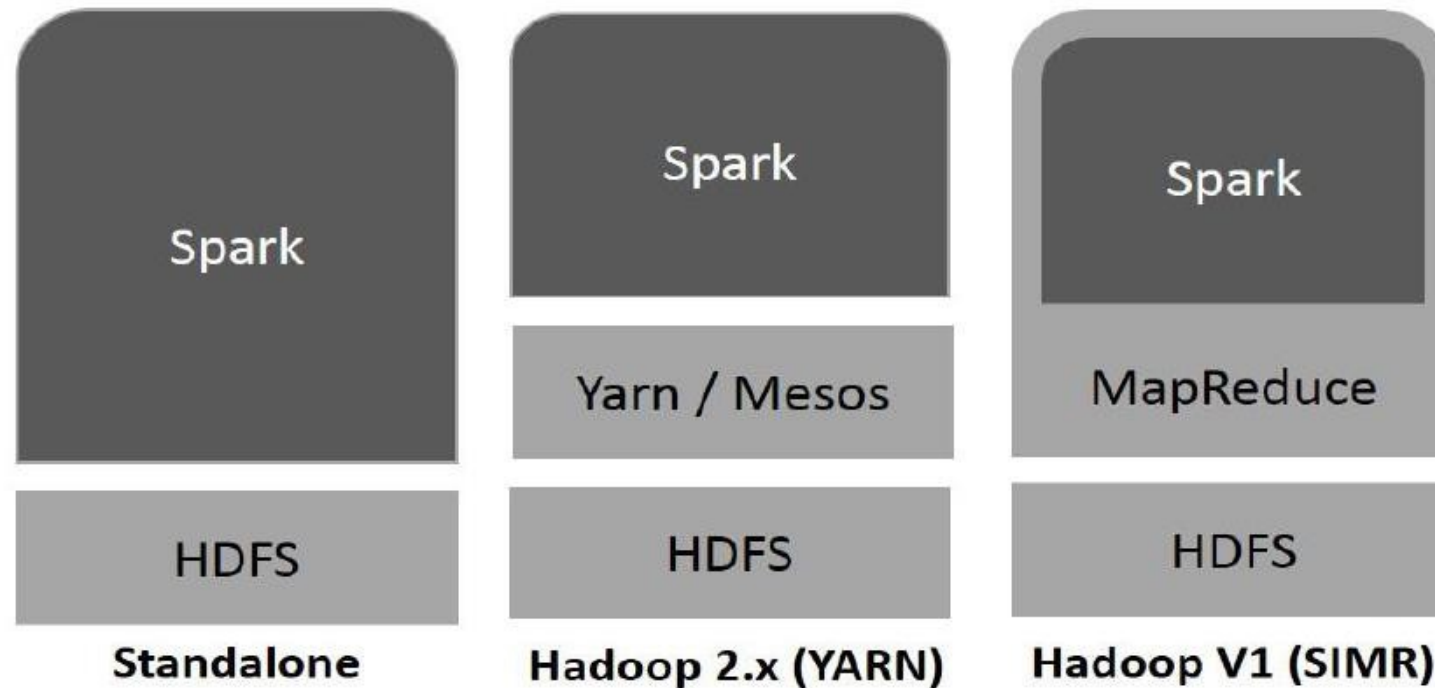
**Spark is not a modified version of Hadoop nor is dependent on Hadoop**

- It has its own cluster management.
- Hadoop is just one of the ways to implement Spark.

**Spark uses Hadoop in two ways**

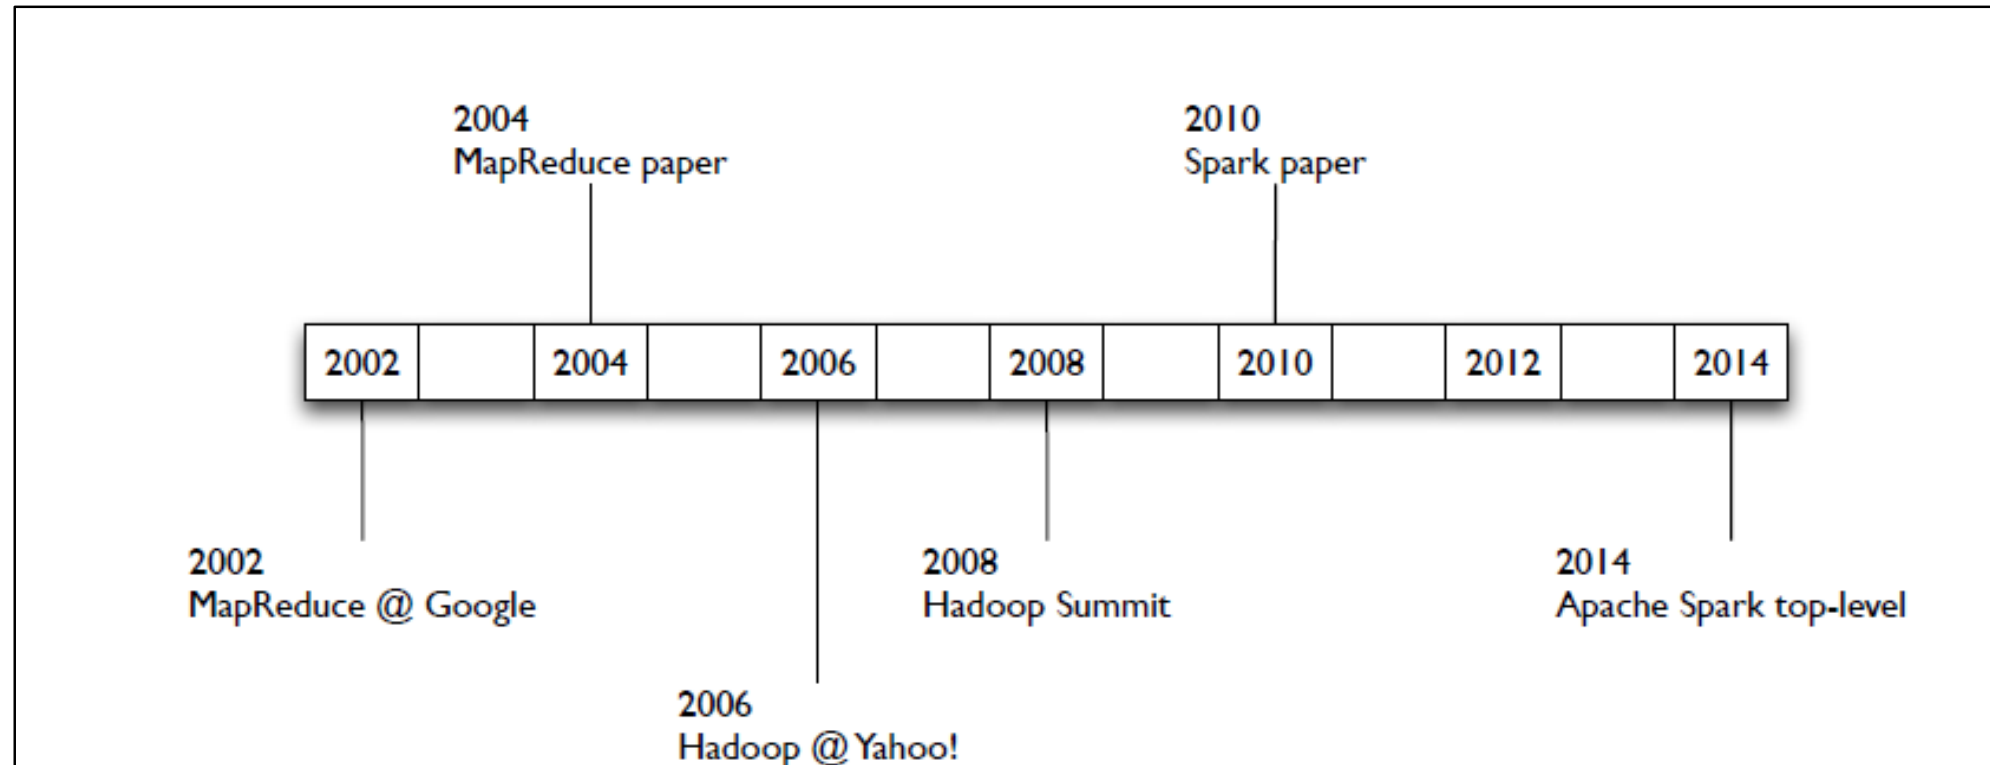
- one is **storage** and
- second is **processing**.

# Spark Built on Hadoop



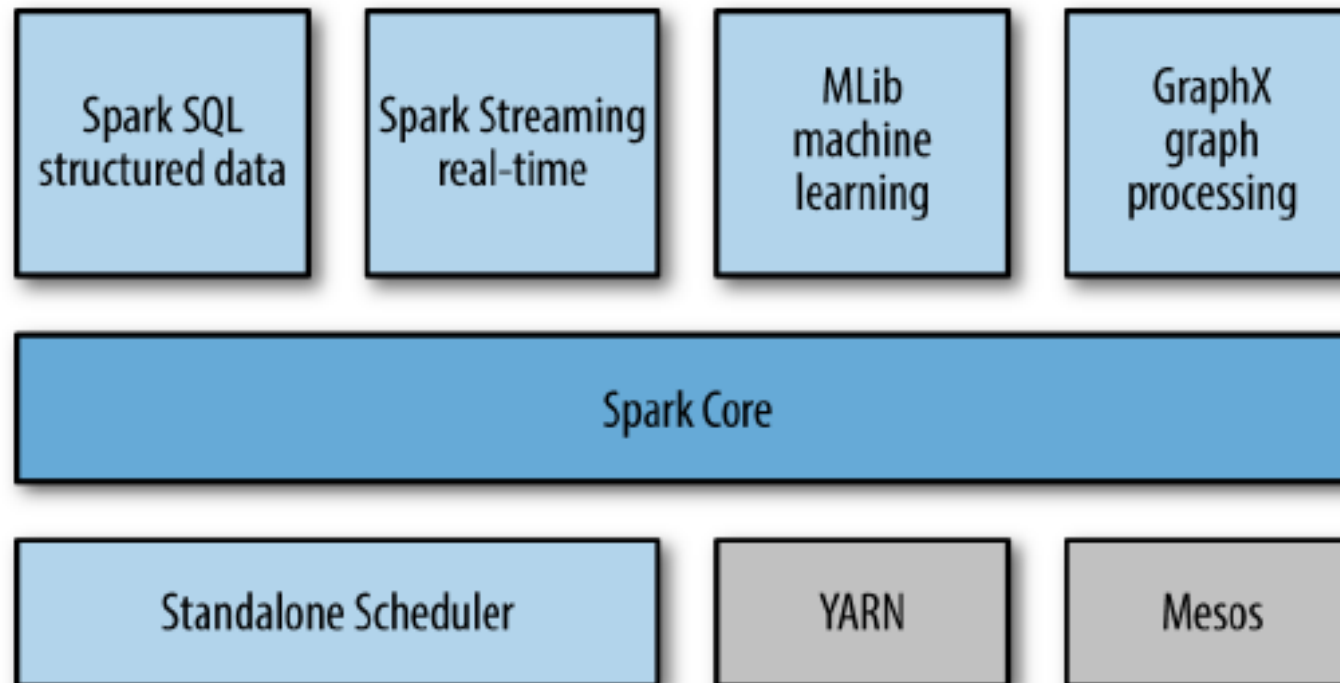
# Brief History

- Developed at UC Berkeley AMPLab in 2009
- Open-sourced in 2010
- Commercial support provided by DataBricks



# Components of Spark

- Spark core is a “computational engine”
  - Deals with scheduling, distributing, and monitoring applications across many worker machines, or a *computing cluster*
- Tight integration helps build applications that combine different processing models



# Resilient Distributed Dataset (RDD)



Fundamental data structure of Spark
Immutable distributed collection of objects
Partitioned collection of records that can be operated on in parallel

Spark use the concept of RDD to achieve faster and efficient MapReduce operations.

Let's design a data processing language "from scratch"!



(Why is MapReduce the way it is?)


When and why they are not so efficient.

# Data-Parallel Dataflow Languages

We have a collection of records,  
want to apply a bunch of transformations  
to compute some result

*Assumptions:* static collection,  
records (not necessarily key-value pairs)

We need per-record processing  
(note, not necessarily key-value pairs)



*Remarks:* Easy to parallelize maps,  
record to “mapper” assignment is an implementation detail



(If we want more than embarrassingly parallel processing)

Map alone isn't enough


We need a way to group partial results  
Intermediate (key, value) pairs

grouping key      partial result


The diagram consists of two labels at the bottom: 'grouping key' on the left and 'partial result' on the right. From 'grouping key', an arrow points diagonally up and to the right towards the word 'key' in the text 'Intermediate (key, value) pairs' above. From 'partial result', an arrow points diagonally up and to the left towards the word 'value' in the same text.

For each key, we can apply some computation


# MapReduce



# MapReduce





# MapReduce Workflows




What's wrong?

# Want MM...?



# Want MRR?



## Iterative Operations on MapReduce




Figure: Iterative operations on MapReduce

Key idea of spark is  
RDD

Data sharing in memory is  
10 to 100 times faster  
than network and Disk.

RDD supports in-  
memory processing

- it stores the state of memory as an object across the jobs
- the object is sharable between those jobs.

## Iterative Operations on Spark RDD

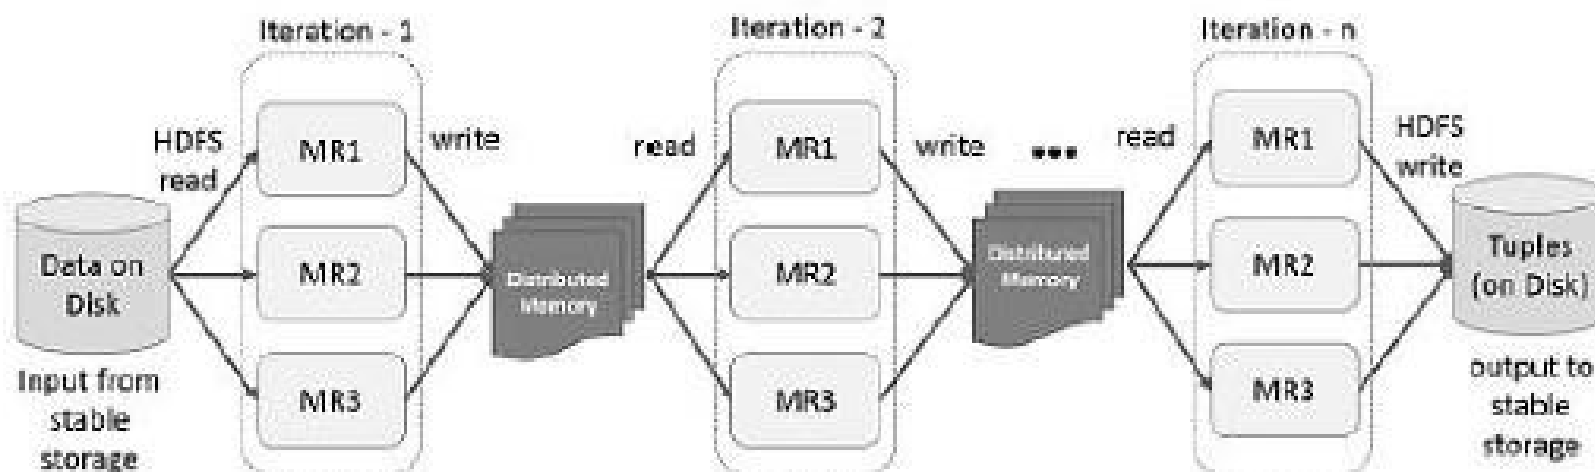
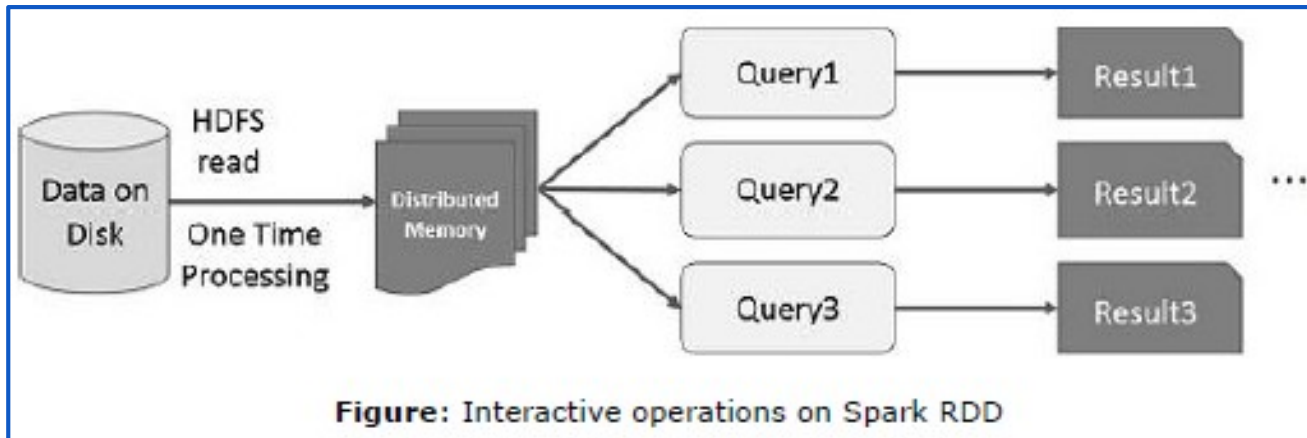
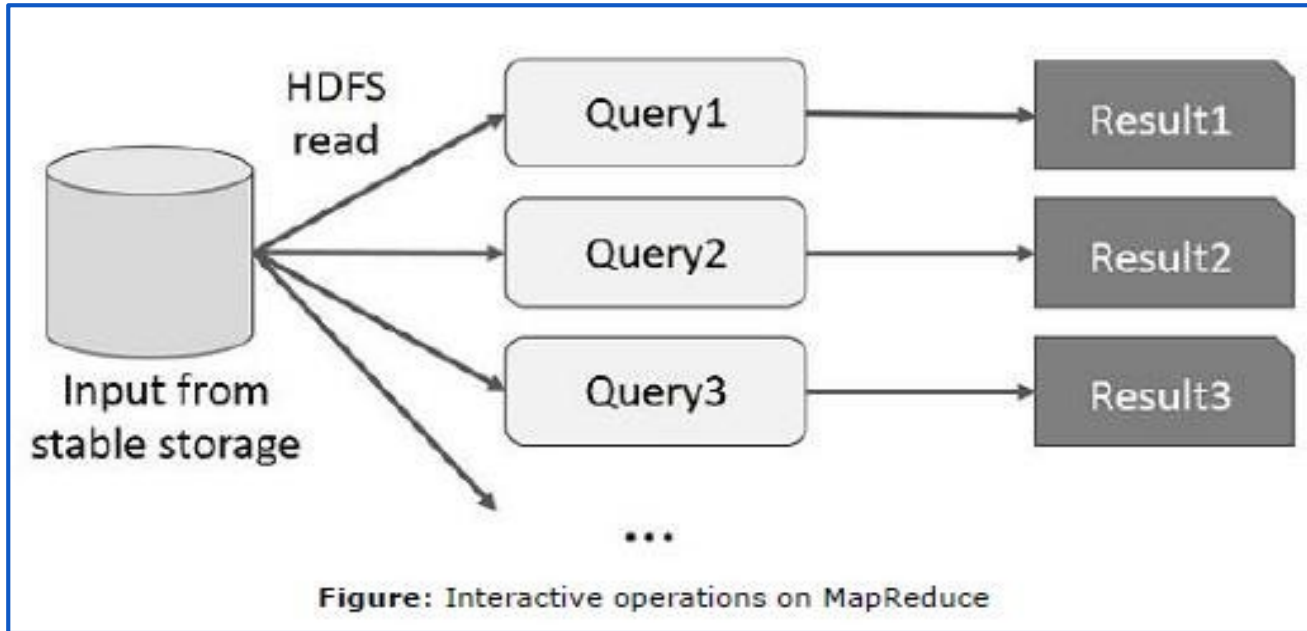


Figure: Iterative operations on Spark RDD

## Interactive Operations



Interactive operations on Spark RDD

For different queries on the same data set keep data in memory

RDD's can be recomputed when an action is run on it

RDD can be persisted in memory, or on disk, or replicated across multiple nodes



# Spark-Resilient Distributed Datasets (RDD)

- An RDD is simply a distributed collection of elements
  - Core abstraction
  - Immutable
  - Partitioned collection of objects spread across a cluster,
    - stored in RAM or on disk
- Achieve fault tolerance through lineage
  - if a partition of an RDD is lost, the RDD has enough information about how it was derived from other RDDs to be able to rebuild just that partition.

# Spark-RDD

- Simple view
  - RDD is collection of data items split into partitions and stored in memory on worker nodes of the cluster
- Complex view
  - RDD is an interface for data transformation
  - RDD refers to the data stored either in persisted store (HDFS, Cassandra, HBase, etc.) or in cache (memory, memory+disks, disk only, etc.) or in another RDD

# Spark-Operations on RDD

- In Spark all work is expressed as either
  - creating new RDDs,
  - transforming existing RDDs, or
  - calling operations on RDDs to compute a result.
- Under the hood, Spark automatically distributes the data contained in RDDs across your cluster and parallelizes the operations you perform on them

# Spark-Transformations


- *Transformations* construct a new RDD from a previous one.
- Example of transformation is filtering data that matches a predicate.

RDD transformations returns pointer to new RDD and allows you to create dependencies between RDDs.

**Transformations in Spark are “lazy”.** They do not compute results immediately.


They just “remember” the operation to be performed and the dataset (e.g., file) to which the operation is to be performed

This design enables Spark to run more efficiently. WHY?




# Spark-Actions

- *Actions* compute a result based on an RDD, and either return it to the driver program or save it to an external storage system (e.g., HDFS).
- Examples of actions include
  - count (which returns the number of elements in the dataset),
  - collect (which returns the elements themselves), and
  - save (which outputs the dataset to a storage system)



# Spark-Transformation vs Actions

- Transformations and actions are different because of the way Spark computes RDDs
  - Spark computes RDDs in a *lazy* fashion—that is, the first time they are used in an action.
- Spark's RDDs are by default recomputed each time you run an action on them.
  - To reuse an RDD in multiple actions, you can ask Spark to *persist* it using **RDD.persist()**



# Create RDD

- Spark provides an interactive shell: a powerful tool to analyze data interactively

Open spark shell

```
$ spark-shell
```

- RDDs can be created from
  - Hadoop Input Formats (such as HDFS files) or
  - Transforming other RDDs

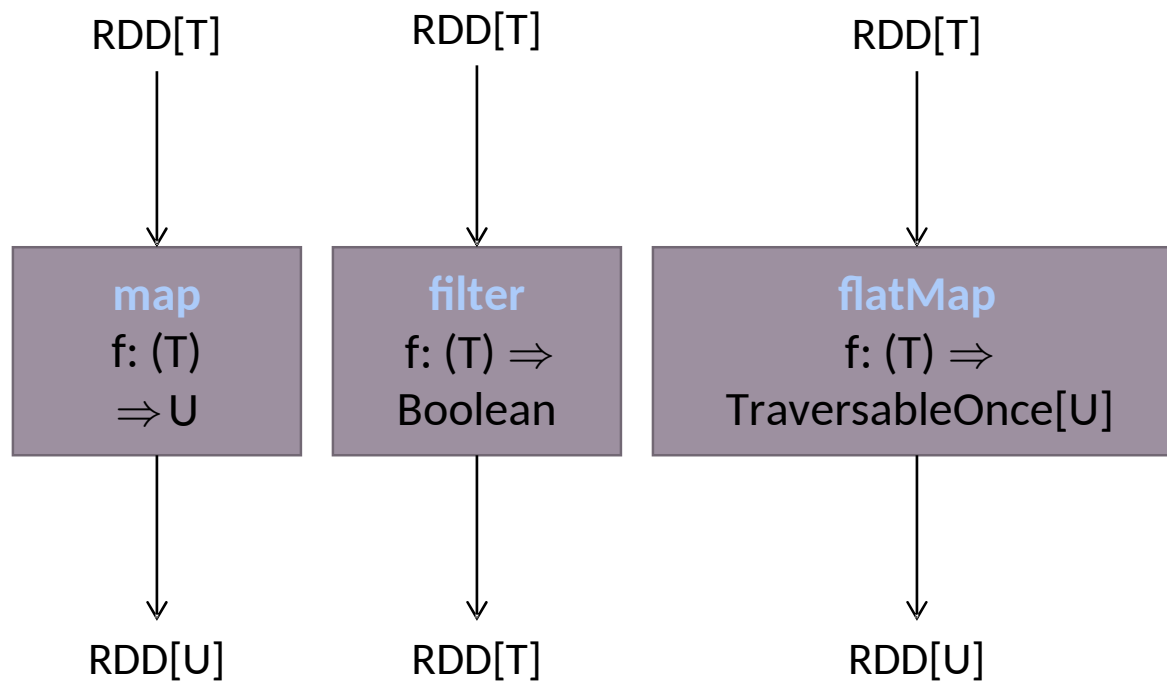
Create Simple RDD

```
scala> val inputfile = sc.textFile("input.txt")
```

```
scala> sc.parallelize (List(1,2,3))
```

[Quick Scala Tutorial](#)

# RDD Transformations



## Map-like Operations

(Not meant to be exhaustive)



# Spark Basic Transformations

- `val rdd = sc.parallelize(List(1,2,3))`
- `val rdd2 = rdd.filter(x=> x%2 == 0)`
  - Pass each element through a function, keep those fulfil the predicate
  - rdd2 will be {2}
- `val inputRDD = sc.textFile("log.txt")`
- `val errorsRDD = inputRDD.filter(line => line.contains("error"))`

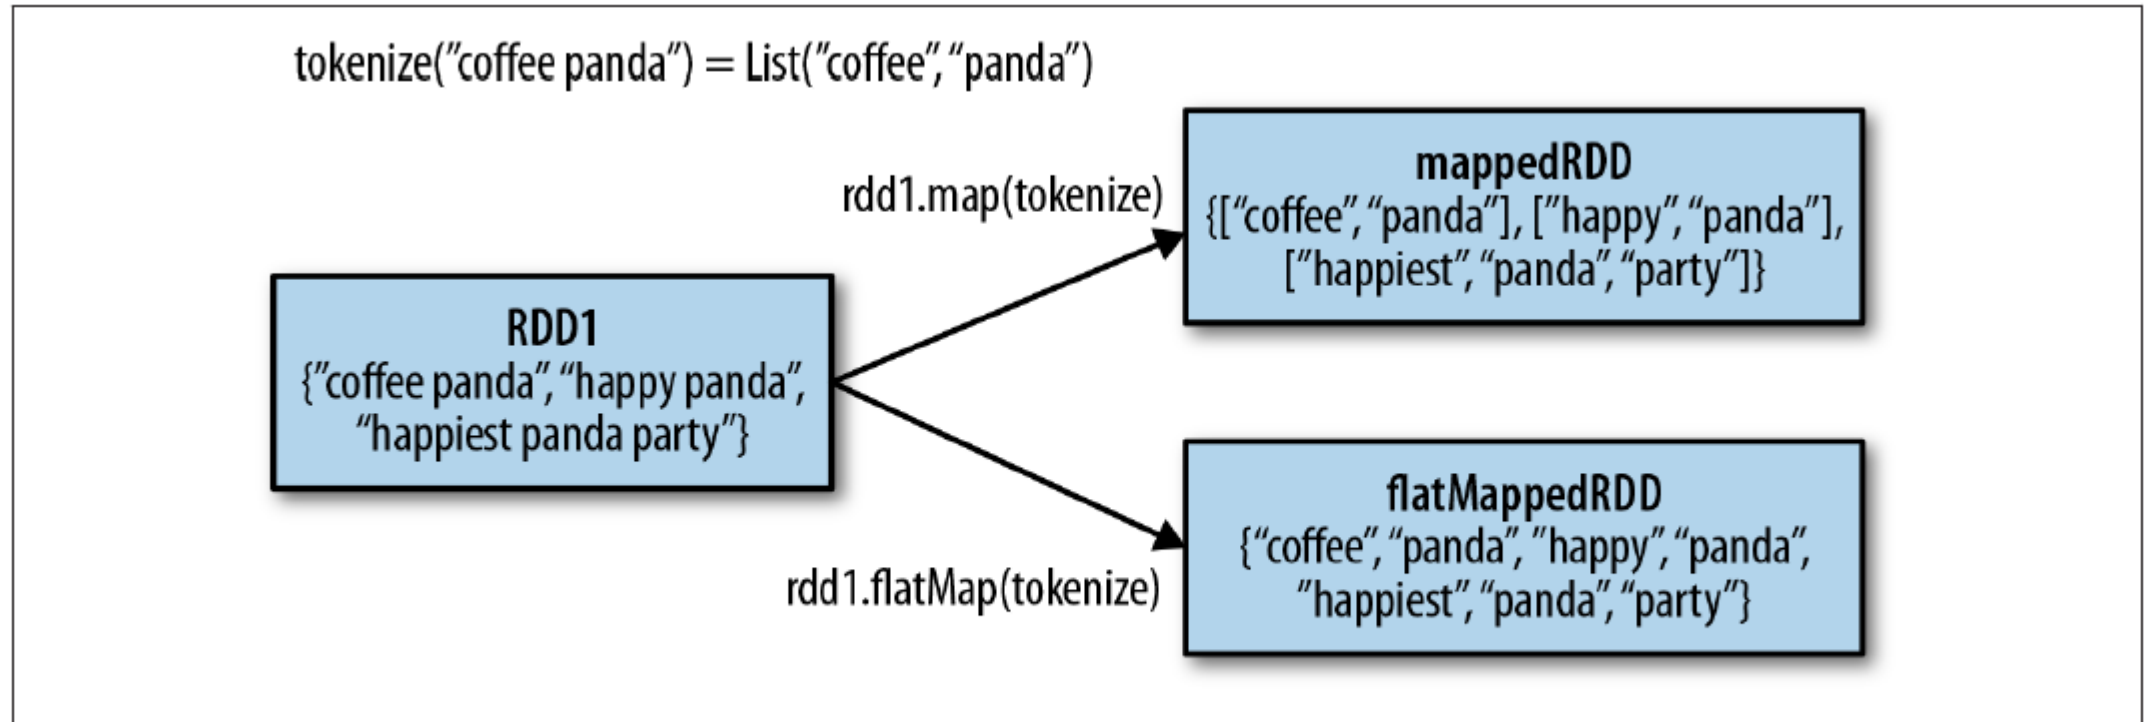
# Spark Basic Transformations

- `val rdd = sc.parallelize(List(1,2,3))`
- `val rdd1 = rdd.map(x=>x*2)`
  - Pass each element through a function, rdd1 will be {1 , 4, 6}
- `val rdd2 = rdd.flatMap(x=> 0 to x)`
  - Map each element to 0 or more other elements and flatten the result into one array,
  - rdd2 will be { 0, 1, 0, 1, 2, 0, 1, 2, 3}
- `val rdd3 = rdd.map(x=> 0 to x)`
  - Map each element to 0 or more other elements and returns a list of arrays
  - rdd3 will be { 0, 1}, {0, 1, 2} , {0, 1, 2, 3}

# Spark Basic Transformations


- `val sData = sc.parallelize(List("Summer days", "Autumn hues", "Winter cold", "Spring flowers"))`
- `val D3 = sData.flatMap(x=>x.split(" "))`
  - **D3 is** {Summer, days, Autumn, hues, Winter, cold, Spring, flowers}
- `val D4 = sData.map(x=>x.split(" "))`
  - **D4 is** Array(Summer, days), Array(Autumn, hues), Array(Winter, cold), Array(Spring, flowers)

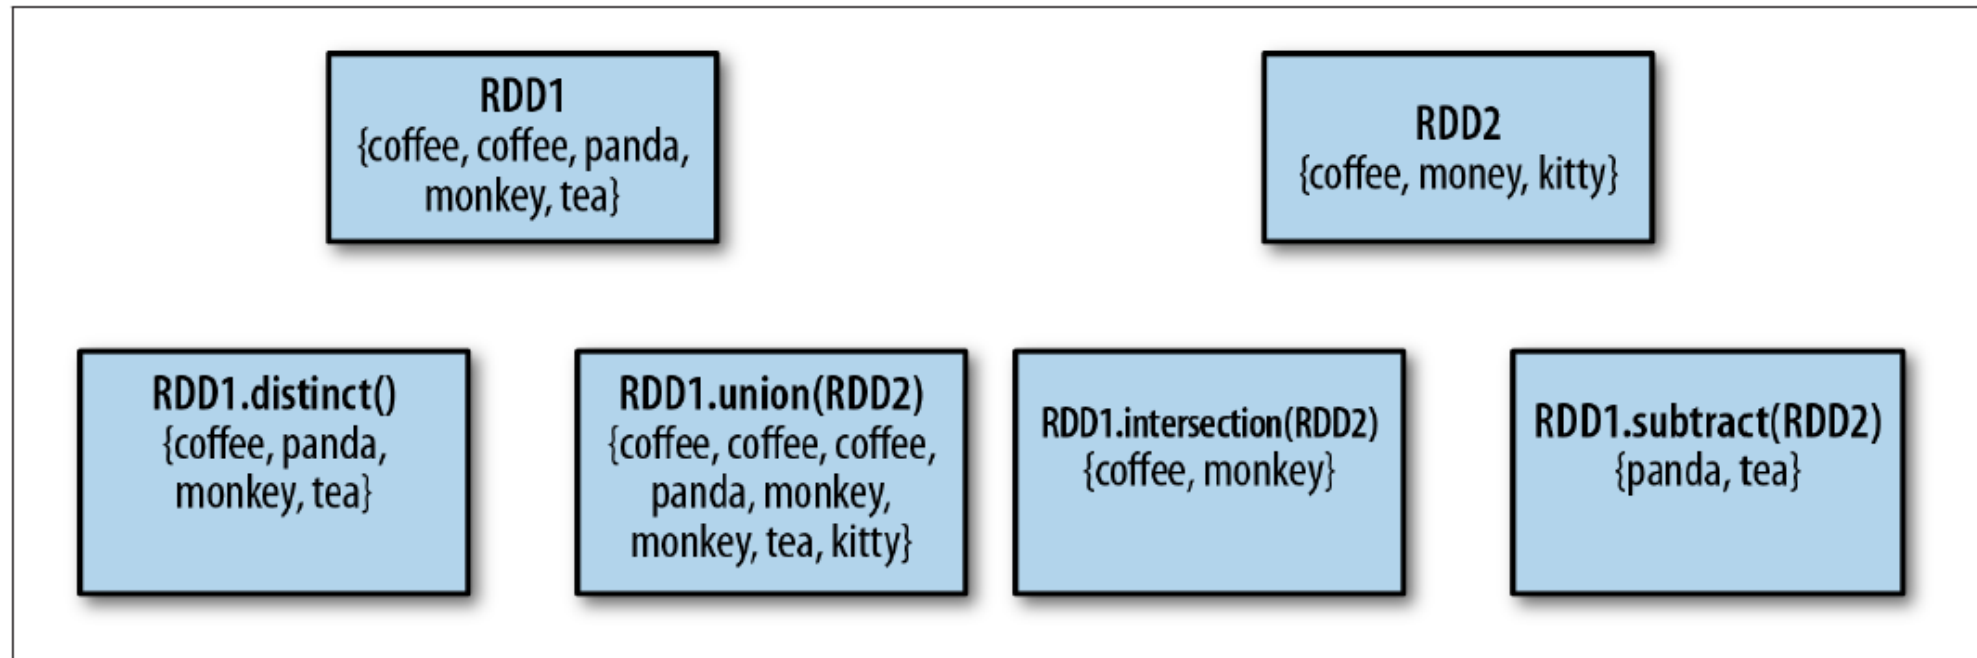
# map() vs flatMap()



*Figure 3-3. Difference between flatMap() and map() on an RDD*

# Set Operations





*Figure 3-4. Some simple set operations*

The set property most frequently missing from UNION is the uniqueness of elements.

The performance of intersection() is much worse than union() since it requires a shuffle over the network to identify common elements.

# Set Operations

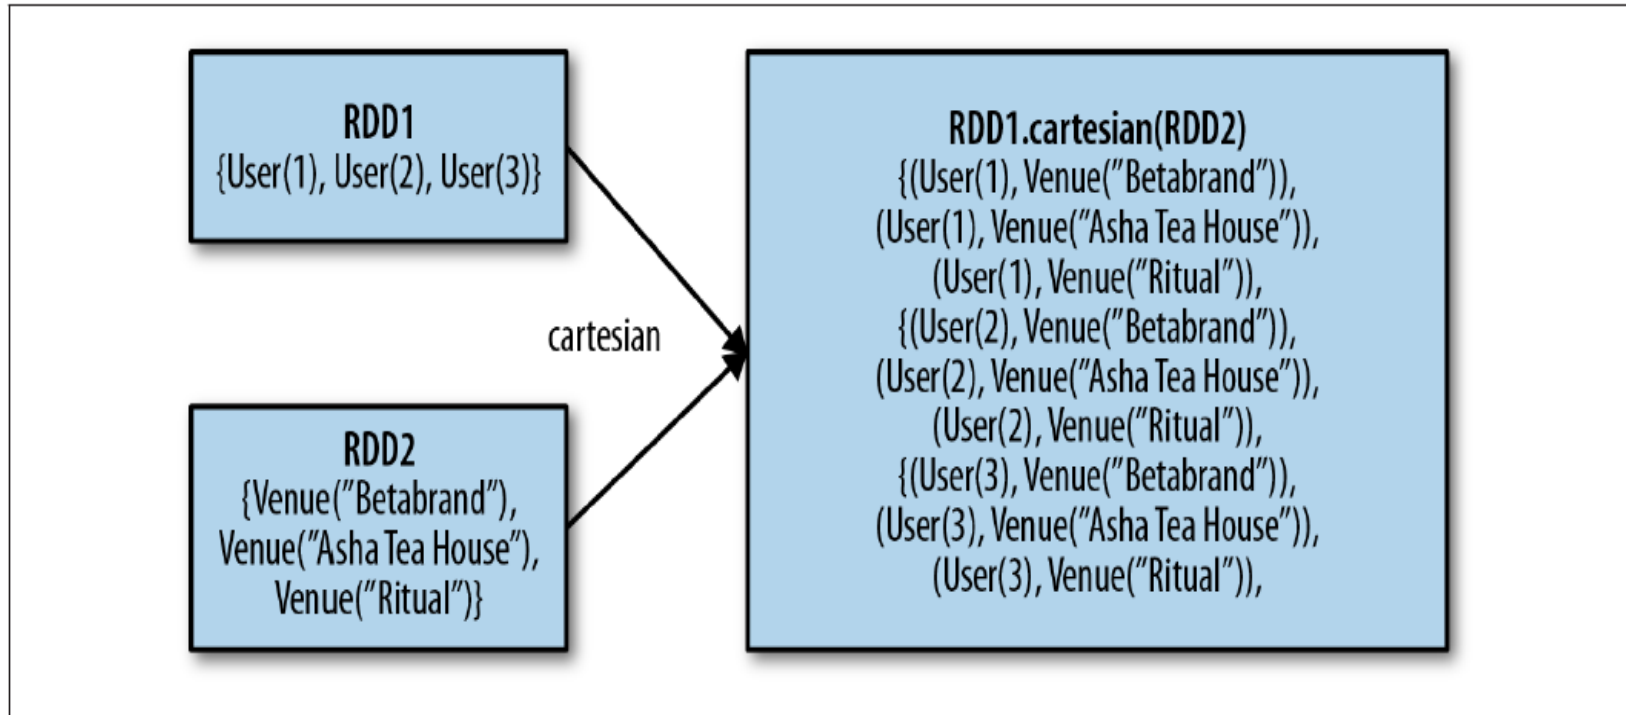




Figure 3-5. Cartesian product between two RDDs



(Not meant to be exhaustive)

# Transformations And Lineage Graph

- Spark keeps track of the set of dependencies between different RDDs, called the *lineage graph*.
- It uses *lineage graph* to compute each RDD on demand and to recover lost data if part of a persistent RDD is lost.



*RDD lineage graph created during log analysis*




# ACTIONS

---

# Actions

- Actions force the evaluation of the transformations required for the RDD they were called on, to produce output.



# Basic Actions

- `val nums = sc.parallelize(List(1,2,3,3,4))`
- `val odd=nums.filter(x => x%2==1)`
- `odd.collect()`
  - retrieve the entire RDD
- `val rddT= nums.take(2)`
- `rddT.foreach(println)`

`foreach(func)` apply the provided function to each element of the RDD.

Keep in mind that your entire dataset must fit in memory on a single machine to use `collect()` on it, so `collect()` shouldn't be used on large datasets.

`take(n)` returns  $n$  elements from the RDD

It's important to note that these operations do not return the elements in the order you might expect.

# Basic Actions

- `val nums = sc.parallelize(List(1,2,3,3,4))`
  1. `nums.saveAsTextFile("filename.txt")`
  2. `nums.foreach(println)`
  3. `nums.count()`
  4. `nums.countByValue()`
    - Number of times each element occurs in the RDD.
    - `{(1, 1),(2, 1),(3, 2),(4,1)}`
  5. `val sum = nums.reduce((x, y) => x + y)`
    - Combine the elements of the RDD together in parallel (e.g., sum).

# Basic Actions

- `val nums = sc.parallelize(List(1,2,3,3,4))`
- `val agg = nums.aggregate((0,0)) (  
 (x, value)=>(x._1 + value, x._2 + 1),  
 (x1, x2)=>(x1._1 + x2._1, x1._2 + x2._2))  
 println(agg)`
- `val average = agg._1 /agg._2.toDouble`
- `println(average)`

Computing sum and count in one action

First provide a function to combine the elements from our RDD with the accumulator.

Next we need a second function to merge two accumulators

# Lazy Evaluation

Transformations  
on RDDs are  
lazily evaluated

- A transformation is not immediately performed.
- Spark records metadata to indicate that this operation has been requested
- **For example:** when we call `sc.textFile()`, the data is not loaded until it is necessary.

WHY?

- Spark uses lazy evaluation to reduce the number of passes it has to take over our data by grouping operations together

# Persistence (Caching)

- `val result = nums.map(x => x*x)`
- `println(result.count())`
- `println(result.collect().mkString(","))`

Especially expensive for  
iterative algorithms

Issue ??

RDD is  
computed  
multiple times

Solution??

Persist the data

# Persistence (Caching)

- `val result = nums.map(x => x*x)`
- `result.persist()`
- `println(result.count())`
- `println(result.collect().mkString(","))`

What happened when a node that has data persisted on it fails ?

Spark will recompute the lost partitions of the data



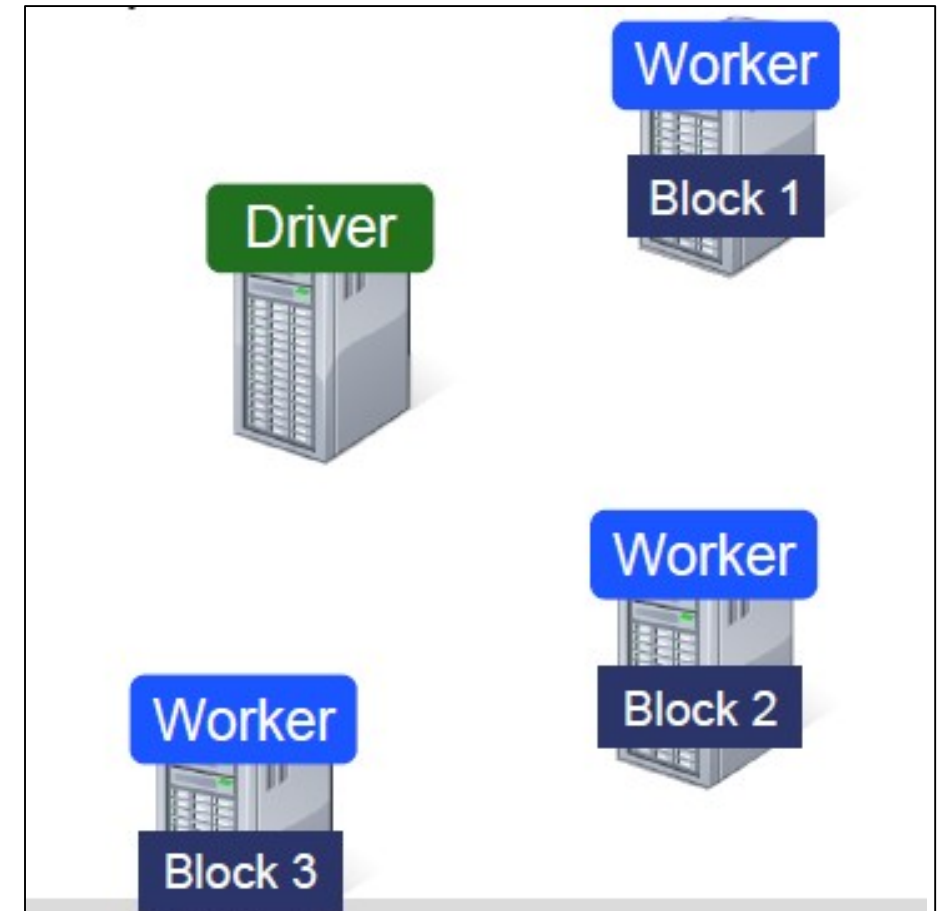
*Table 3-6. Persistence levels from `org.apache.spark.storage.StorageLevel` and `pyspark.StorageLevel`; if desired we can replicate the data on two machines by adding `_2` to the end of the storage level*

Level	Space used	CPU time	In memory	On disk	Comments
MEMORY_ONLY	High	Low	Y	N	
MEMORY_ONLY_SER	Low	High	Y	N	
MEMORY_AND_DISK	High	Medium	Some	Some	Spills to disk if there is too much data to fit in memory.
MEMORY_AND_DISK_SER	Low	High	Some	Some	Spills to disk if there is too much data to fit in memory. Stores serialized representation in memory.
DISK_ONLY	Low	High	N	Y	

```
import org.apache.spark.storage.StorageLevel
result.persist(StorageLevel.DISK_ONLY)
```

# Example: Log Mining

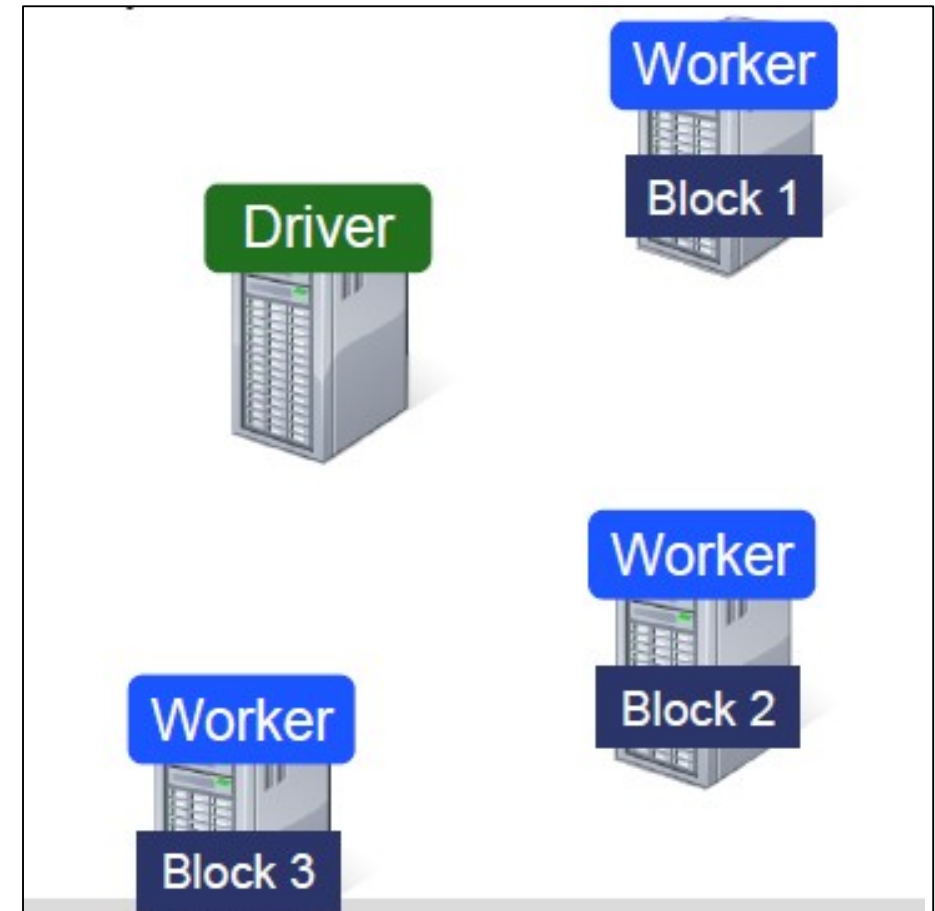
- Load error messages from a log into memory and search for various patterns
- Data is separated into partitions on different machines



# Example: Log Mining

*Load error messages from a log into memory and search for various patterns*

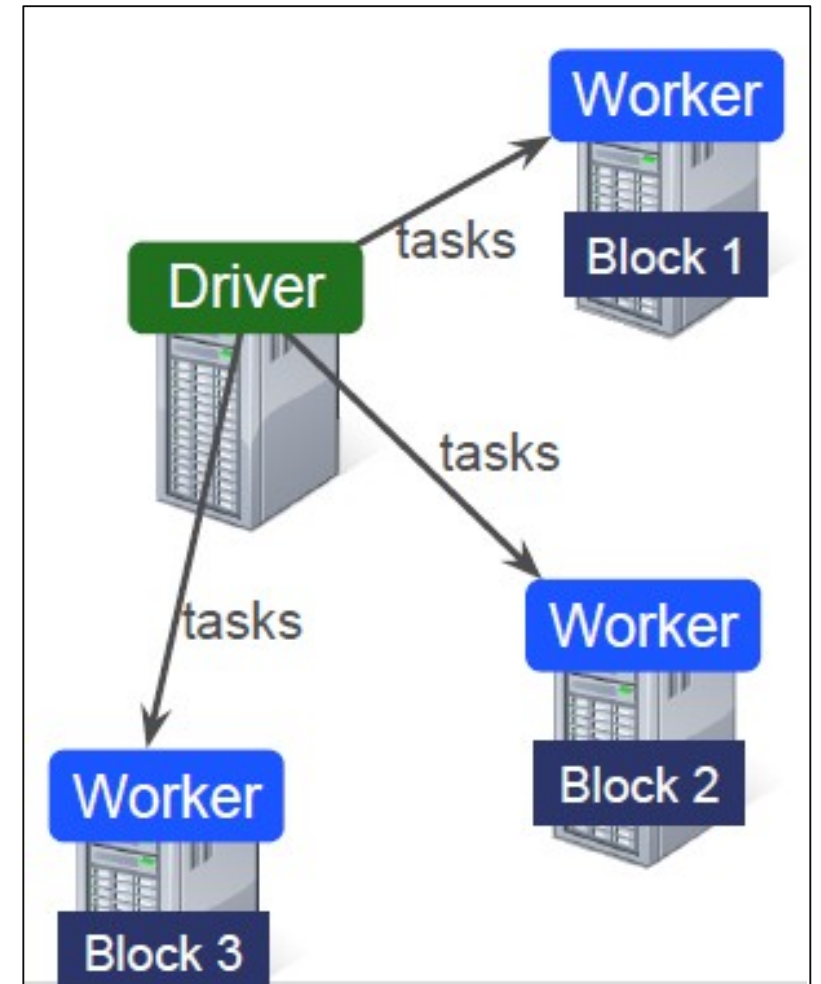
- Suppose that we wish to count the lines containing errors in a large log file stored in HDFS.
  - `val file = spark.textFile("hdfs://...")`
  - `val errs = file.filter(_.contains("ERROR"))`
  - `val cachedErrs = errs.cache()`
  - `val ones = cachedErrs.map(_ => 1)`
  - **`val count = ones.reduce(_+_)`**
- `cachedErrs.filter(_.contains("mysql")).count()`



# Example: Log Mining

*Load error messages from a log into memory and search for various patterns*

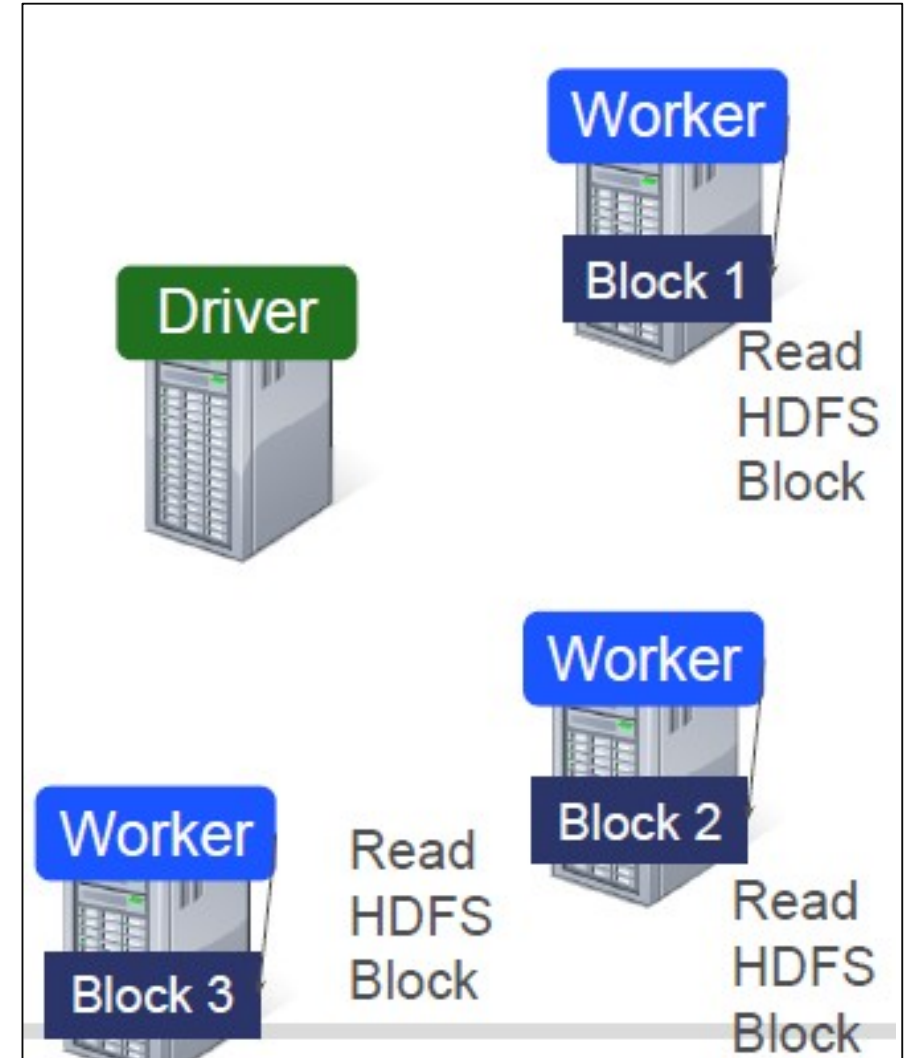
- Suppose that we wish to count the lines containing errors in a large log file stored in HDFS.
  - **val file = spark.textFile("hdfs://...")**
  - **val errs = file.filter(\_.contains("ERROR"))**
  - **val cachedErrs = errs.cache()**
  - **val ones = cachedErrs.map(\_ => 1)**
  - **val count = ones.reduce(\_+\_)**
  - **cachedErrs.filter(\_.contains("mysql")).count()**



# Example: Log Mining

*Load error messages from a log into memory and search for various patterns*

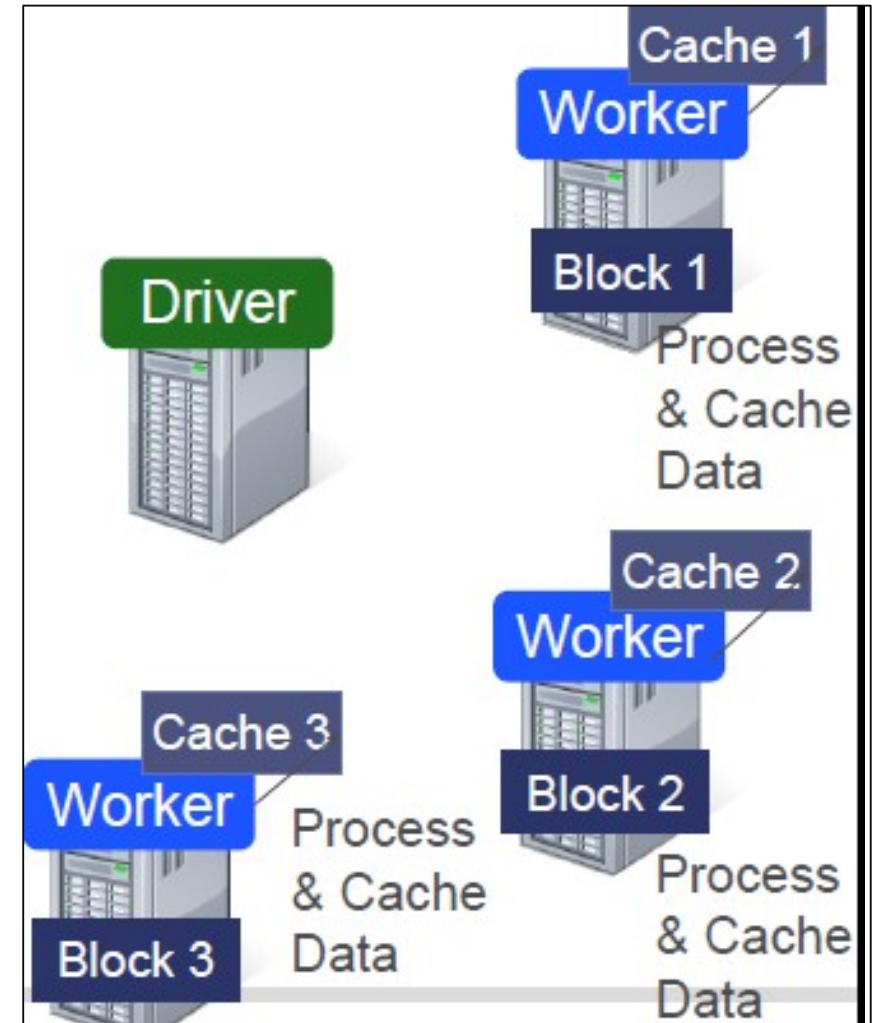
- Suppose that we wish to count the lines containing errors in a large log file stored in HDFS.
  - `val file = spark.textFile("hdfs://...")`
  - `val errs = file.filter(_.contains("ERROR"))`
  - `val cachedErrs = errs.cache()`
  - `val ones = cachedErrs.map(_ => 1)`
  - `val count = ones.reduce(_+_)`
  - `cachedErrs.filter(_.contains("mysql")).count()`



# Example: Log Mining

*Load error messages from a log into memory and search for various patterns*

- Suppose that we wish to count the lines containing errors in a large log file stored in HDFS.
  - `val file = spark.textFile("hdfs://...")`
  - `val errs = file.filter(_.contains("ERROR"))`
  - `val cachedErrs = errs.cache()`
  - `val ones = cachedErrs.map(_ => 1)`
  - `val count = ones.reduce(_+_)`
  - `cachedErrs.filter(_.contains("mysql")).count()`

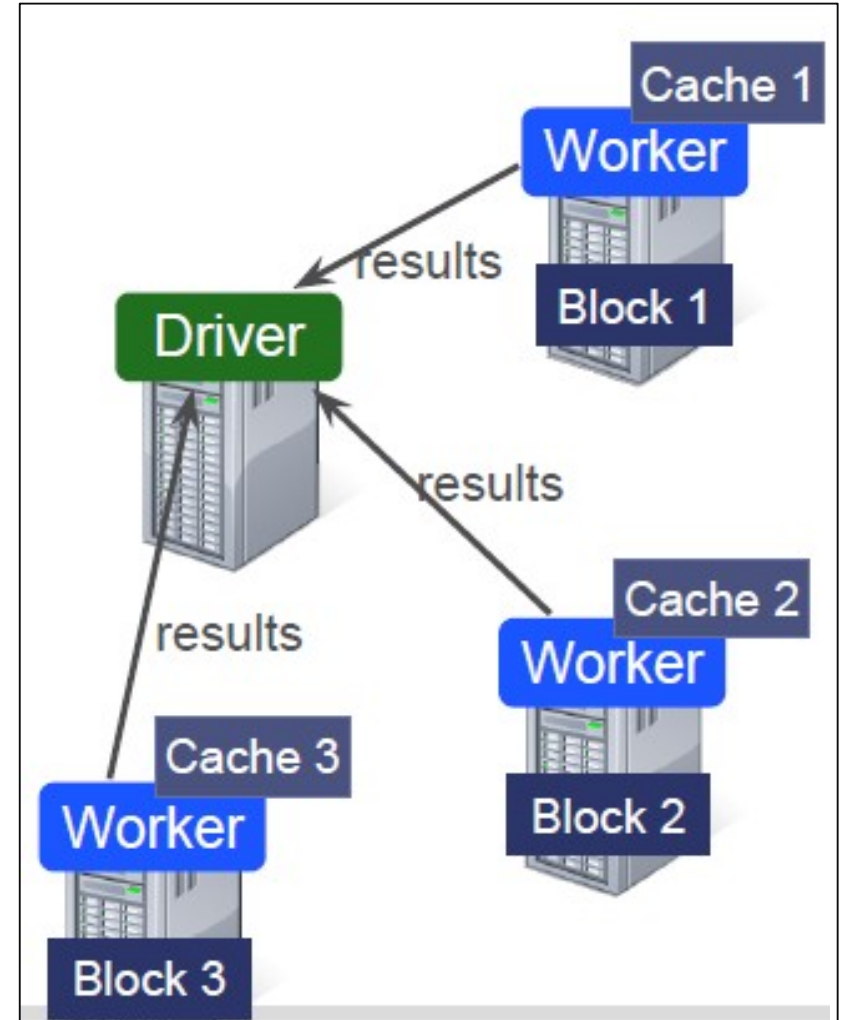




# Example: Log Mining

*Load error messages from a log into memory and search for various patterns*

- Suppose that we wish to count the lines containing errors in a large log file stored in HDFS.
  - `val file = spark.textFile("hdfs://...")`
  - `val errs = file.filter(_.contains("ERROR"))`
  - `val cachedErrs = errs.cache()`
  - `val ones = cachedErrs.map(_ => 1)`
  - **`val count = ones.reduce(_+_)`**
  - `cachedErrs.filter(_.contains("mysql")).count()`



# Example: Log Mining

*Load error messages from a log into memory and search for various patterns*

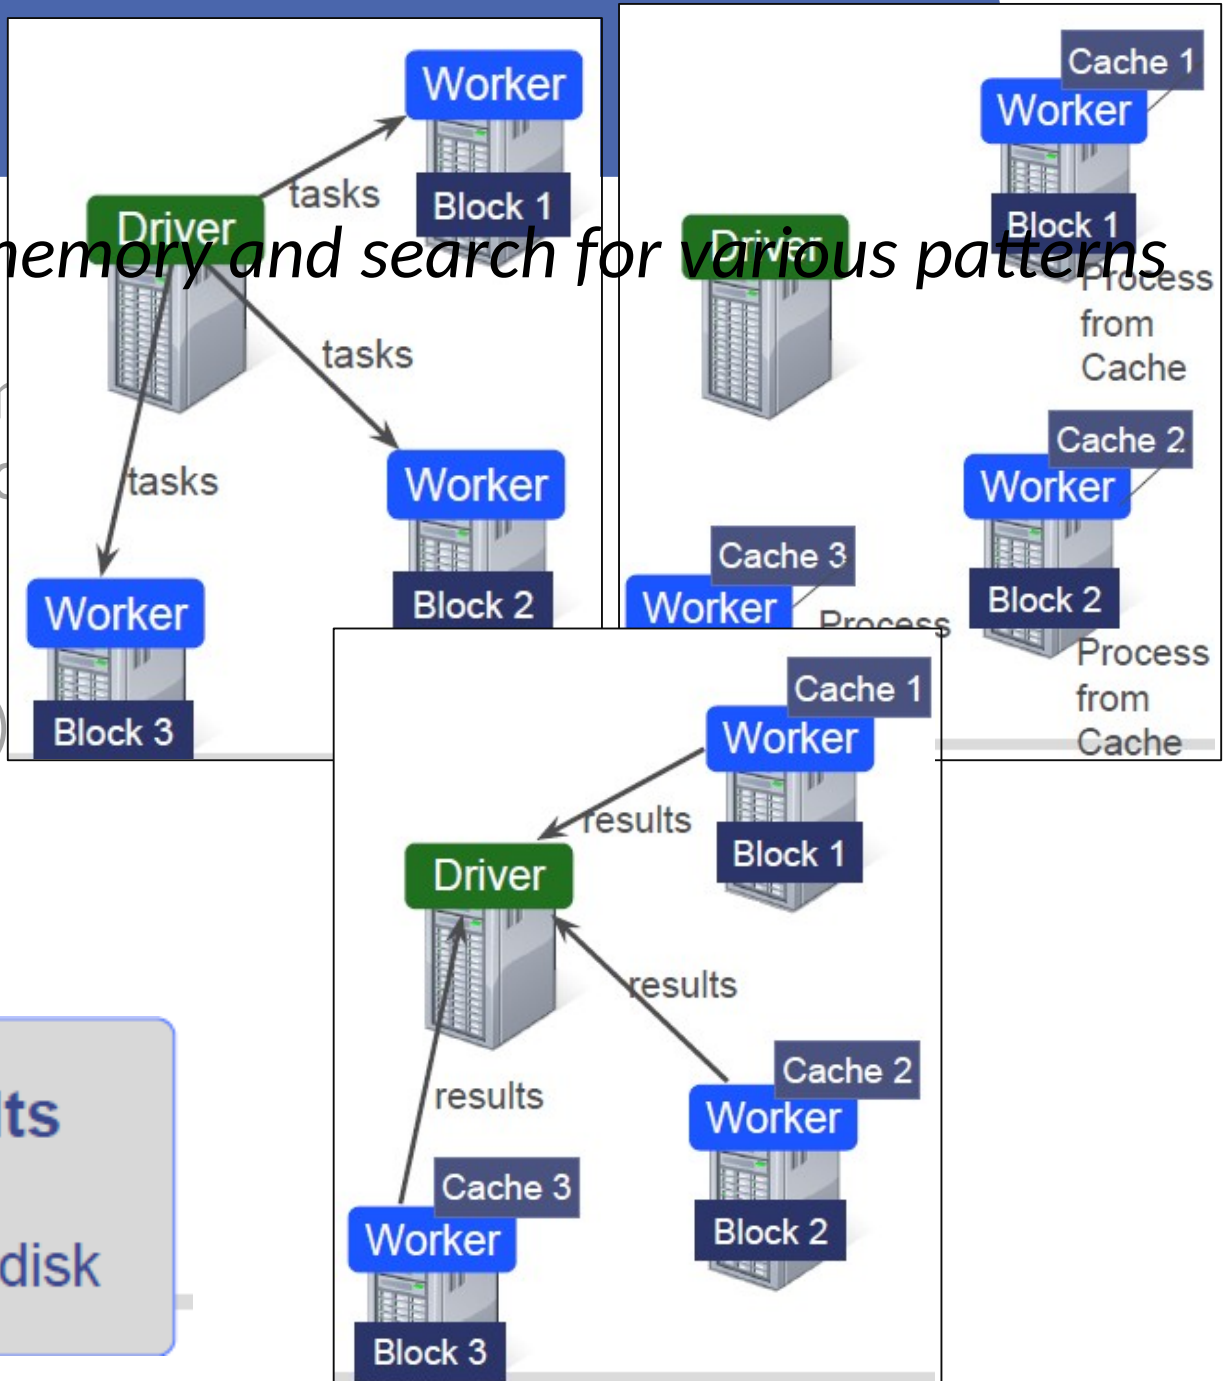
- Suppose that we wish to count the lines containing errors in a large log file stored in HDFS.

- `val file = spark.textFile("hdfs://...")`
- `val errs = file.filter(_.contains("ERROR"))`
- `val cachedErrs = errs.cache()`
- `val ones = cachedErrs.map(_ => 1)`
- `val count = ones.reduce(_+_)`

**Cache your data → Faster Results**

**1 TB of log data data**

- 5-7 sec from cache vs. 170s for on-disk





# Example: Log Mining

- Suppose that we wish to count the lines containing errors in a large log file stored in HDFS.
  - `val file = spark.textFile("hdfs://...")`
  - `val errs = file.filter(_.contains("ERROR"))`
  - `val cachedErrs = errs.cache()`
  - `val ones = cachedErrs.map(_ => 1)`
  - `val count = ones.reduce(_+_)`
  - `cachedErrs.filter(_.contains("mysql")).count()`

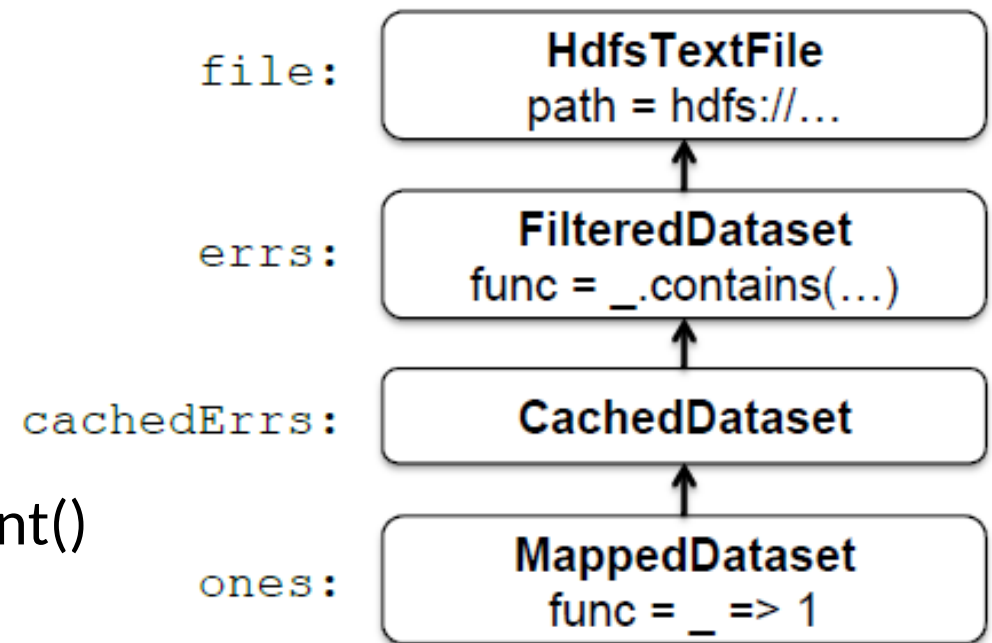


Figure 1: Lineage chain for the distributed dataset objects