

Internet of Things

IO 404 I


Transport Layer

# Transport Layer

- Transport protocols reside on top of IP
- Applications do not use IP directly,
  - but use the transport protocols to communicate with each other.
- ❖ In IP protocol stack, most widely used transport protocols
  - User Datagram Protocol (UDP), and
  - Transport Control Protocol (TCP).

# Transport services and protocols

- ❖ provide *logical communication* between app processes running on different hosts
- ❖ transport protocols run in end systems
  - **send side:** breaks app messages into *segments*, passes to network layer
  - **rcv side:** reassembles segments into messages, passes to app layer
- ❖ more than one transport protocol available to apps
  - Internet: TCP and UDP



# Transport Layer

- **UDP** is a best-effort delivery service, which does not add much on top of IP, whereas
- **TCP** is a reliable byte stream that adds a connection abstraction on top of the connectionless IP.
- Although there have been several other transport protocols defined,
  - such as SCTP [229] and DCCP [152],
  - they have as yet to be adopted by the mainstream.

# Transport Layer

- Usually, basic unit of transportation is called a Packet
- data from higher layers are transported in these packets
- In UDP, the basic unit of transportation is called a user datagram (or UDP segment)
- In TCP, basic unit of transportation is called a segment (TCP segment)

# UDP: User Datagram Protocol [RFC 768]

## **Best effort delivery service:**

- ❖ As the underlying IP network does its best to deliver the datagram,
  - but does not guarantee delivery of the datagrams at the destination [may loss]
  - Does not guarantee that the datagrams are delivered in the same order as they were sent [can be delivered out of order]

# UDP: User Datagram Protocol [RFC 768]

## Connectionless

- ❖ no handshaking between UDP sender, receiver
- ❖ each UDP segment handled independently of others

## Uses

- ❖ streaming multimedia apps (loss tolerant, rate sensitive):  
real time audio/video
- ❖ DNS look ups

## reliable transfer over UDP

- ❖ add reliability at application layer
- ❖ application-specific error recovery!

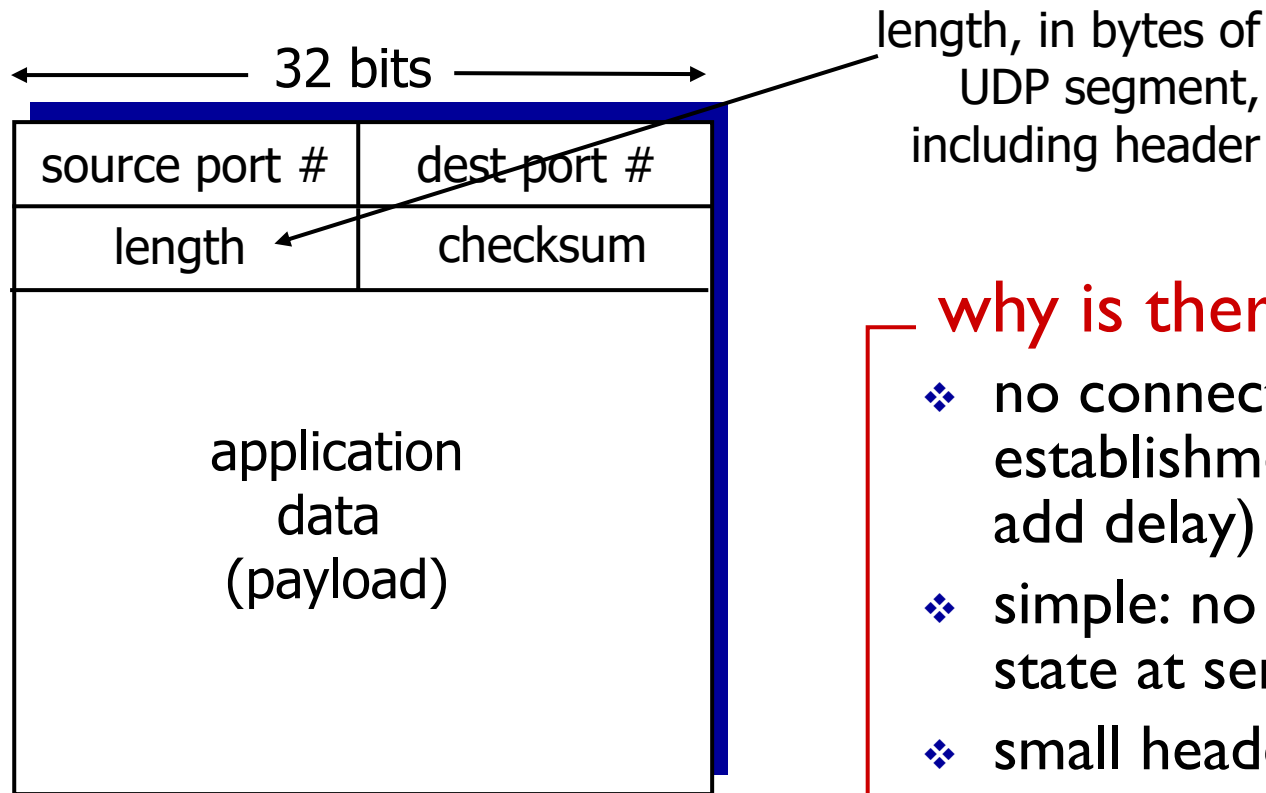
**In the context of smart object networks,**

**UDP** [due to its simplicity and lightweight nature]

- **is an exciting choice for quick transportation of sensor data**



# UDP: segment header



UDP segment format

## why is there a UDP?

- ❖ no connection establishment (which can add delay)
- ❖ simple: no connection state at sender, receiver
- ❖ small header size
- ❖ no congestion control: UDP can blast away as fast as desired

# UDP checksum

*Goal:* detect “errors” (e.g., flipped bits) in transmitted segment

## sender:

- ❖ treat segment contents, including header fields, as sequence of 16-bit integers
- ❖ **checksum:** addition (one's complement sum) of segment contents
- ❖ sender puts checksum value into UDP checksum field

## receiver:

- ❖ compute checksum of received segment
- ❖ check if computed checksum equals checksum field value:
  - NO - error detected
  - YES - no error detected.  
*But maybe errors nonetheless? More later*  
....

# Internet checksum: example

example: add two 16-bit integers

	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
<hr/>																
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1
<hr/>																
sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

*Note:* when adding numbers, a carryout from the most significant bit needs to be added to the result

# TCP: Overview

RFCs: 793, 1122, 1323, 2018, 2581

---

- ❖ **point-to-point:**
  - one sender, one receiver
- ❖ **reliable, in-order *byte stream*:**
  - no “message boundaries”
- ❖ **pipelined:**
  - TCP congestion and flow control set window size
- ❖ **full duplex data:**
  - bi-directional data flow in same connection
  - MSS: maximum segment size
- ❖ **connection-oriented:**
  - handshaking (exchange of control msgs) initializes sender, receiver state before data exchange
- ❖ **flow controlled:**
  - sender will not overwhelm receiver

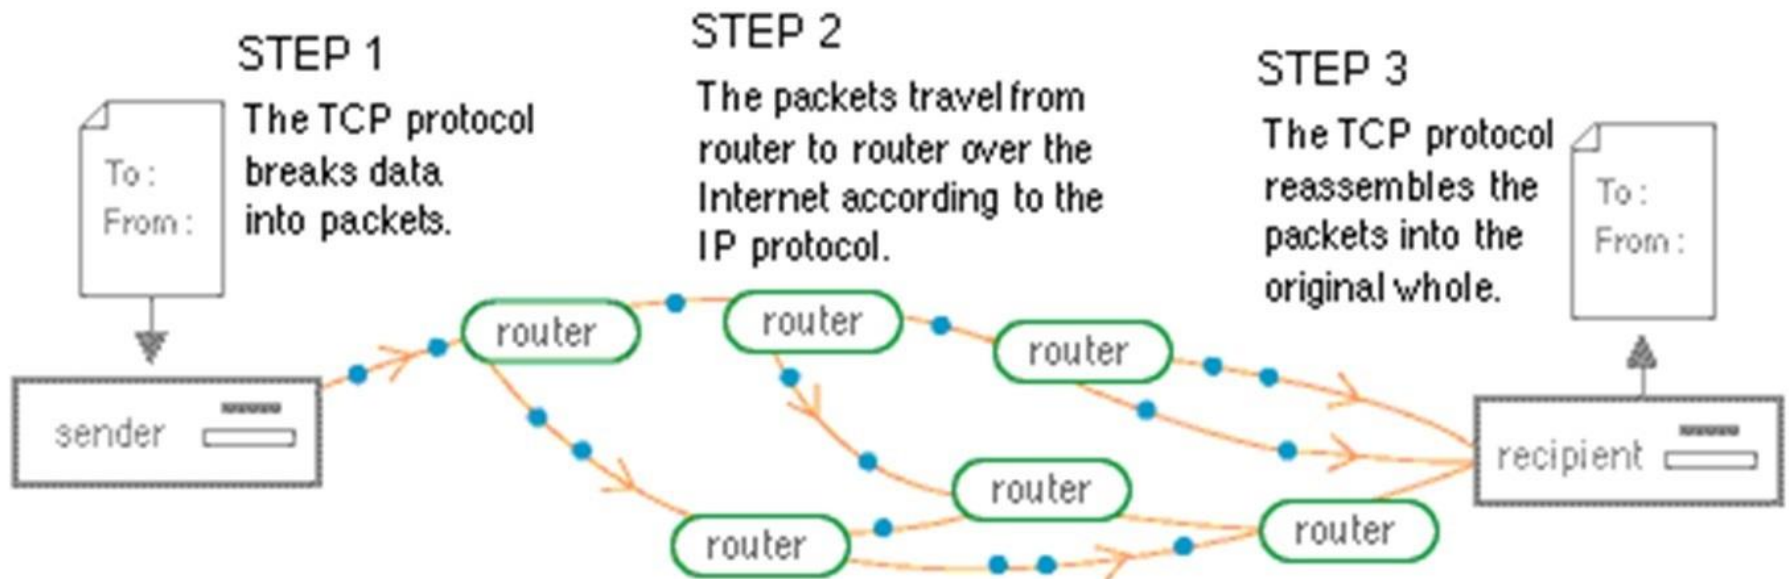
# Transmission Control Protocol (TCP)

- A standard that defines how to establish and maintain a conversation so that application layer protocols can share/exchange data
- provides a reliable byte stream on top of the best-effort packet service provided by the IP layer
- It is based on client-server model
  - Client gets a service from a server, i.e., a web page
- Since IoT devices offer a varying traffic pattern
  - So TCP is not suitable for IoT applications
- ❖ **uIP**, an implementation of TCP for memory-constrained smart objects.

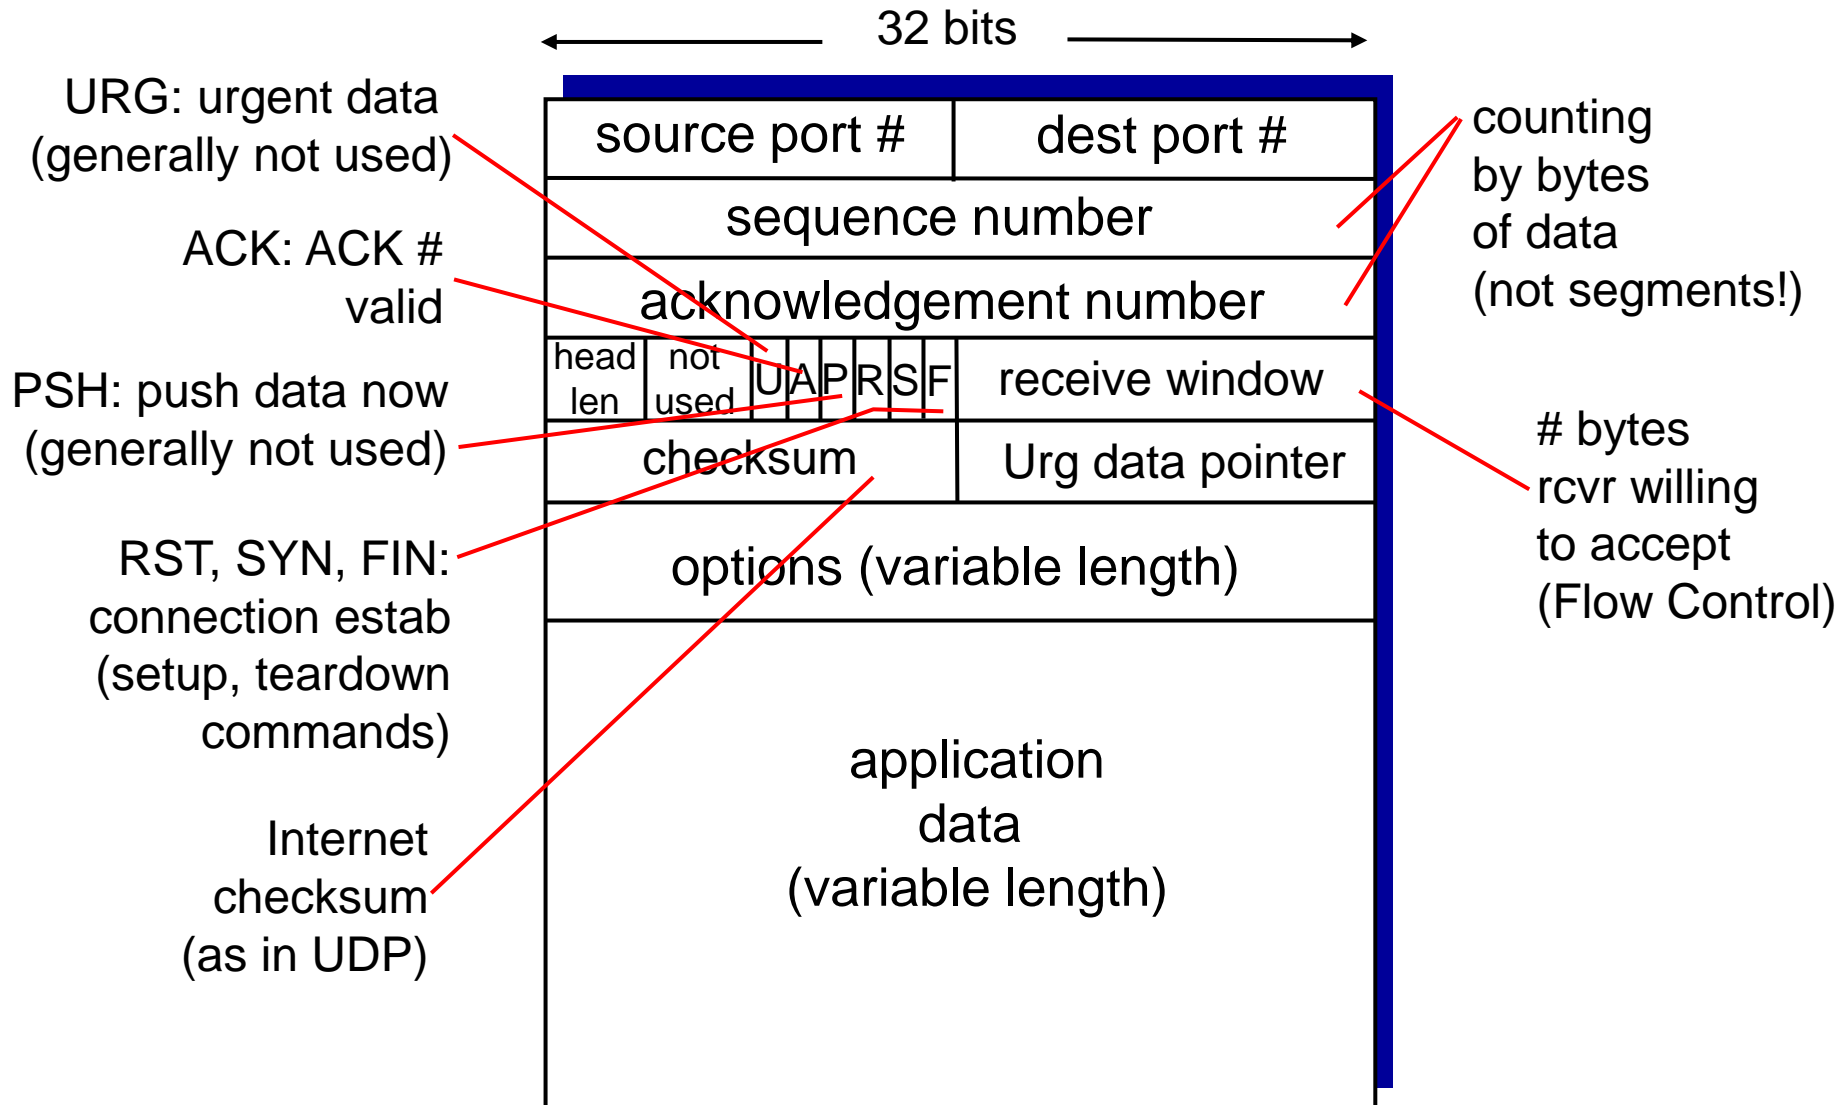
# Transmission Control Protocol (TCP)

- **Reliability** is achieved by **buffering** data combined with positive **ACKs** and **retransmissions**
- Before any data are transported, the two connection end points must explicitly **set up a connection**.
- A connection is identified by the IP addresses and TCP port numbers of the end points
- Many application layer protocols are defined over TCP,
  - such as HTTP (Web), SMTP (e-mail), and XMPP (instant messaging).

# Transmission Control Protocol (TCP)



# TCP segment structure





# TCP seq. numbers, ACKs

## sequence numbers:

- byte stream “number” of the first byte in segment’s data

## acknowledgements:


- seq # of next byte expected from other side
- cumulative ACK

**Q:** how receiver handles out-of-order segments

- A:** TCP spec doesn’t say,  
- up to implementor

outgoing segment from sender


source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer



incoming segment to sender

source port #	dest port #
sequence number	
acknowledgement number	
	<b>A</b> rwnd
checksum	urg pointer

# A simple Telnet Application over TCP



# TCP round trip time, timeout

Q: how to set TCP timeout value?

- ❖ longer than RTT
  - but RTT varies
- ❖ *too short*: premature timeout, unnecessary retransmissions
- ❖ *too long*: slow reaction to segment loss

Q: how to estimate RTT?

- ❖ **SampleRTT**: measured time from segment transmission until ACK receipt
  - ignore retransmissions
- ❖ **SampleRTT** will vary, want estimated RTT “smoother”
  - average several *recent* measurements, not just current **SampleRTT**

# TCP reliable data transfer

❖ TCP creates rdt service on top of IP's unreliable service

- pipelined segments
- cumulative acks
- single retransmission timer

❖ retransmissions triggered by:

- timeout events
- duplicate acks

let's initially consider simplified TCP sender:

- ignore duplicate acks
- ignore flow control, congestion control

# TCP sender events:

## *data rcvd from app:*

- ❖ create segment with seq #
- ❖ seq # is byte-stream number of first data byte in segment
- ❖ start timer if not already running
  - think of timer as for oldest unacked segment
  - expiration interval: `TimeoutInterval`


## *timeout:*

- ❖ retransmit segment that caused timeout
- ❖ restart timer

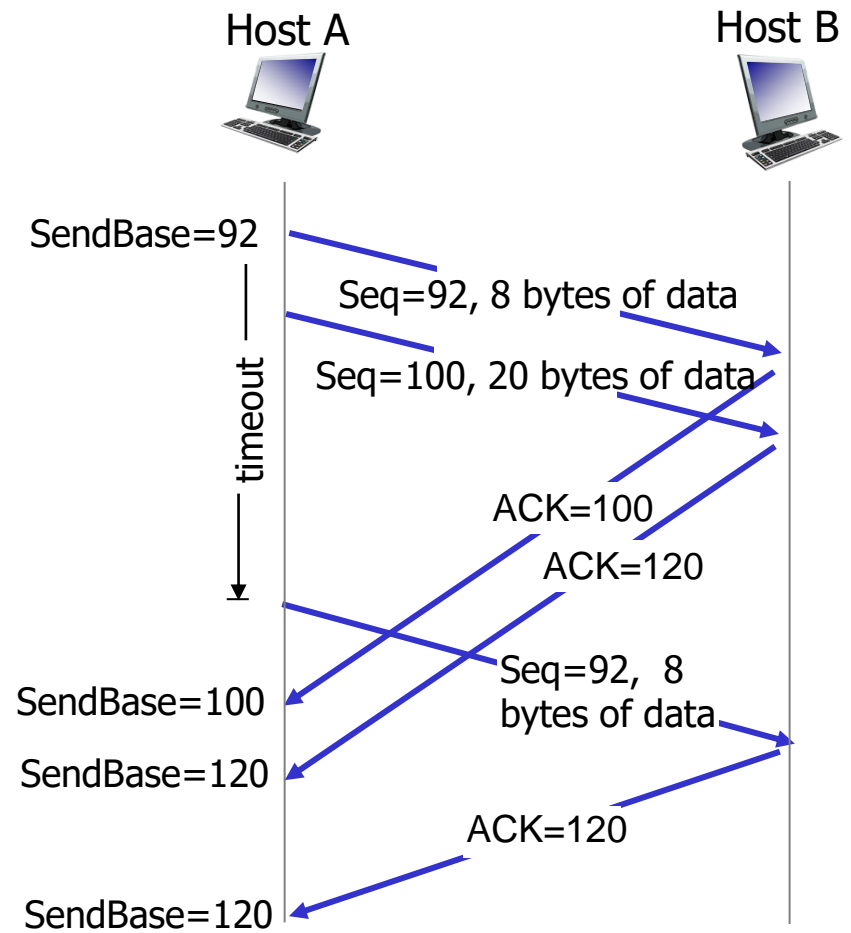
## *ack rcvd:*

- ❖ if ack acknowledges previously unacked segments
  - update what is known to be ACKed
  - start timer if there are still unacked segments

# TCP: retransmission scenarios

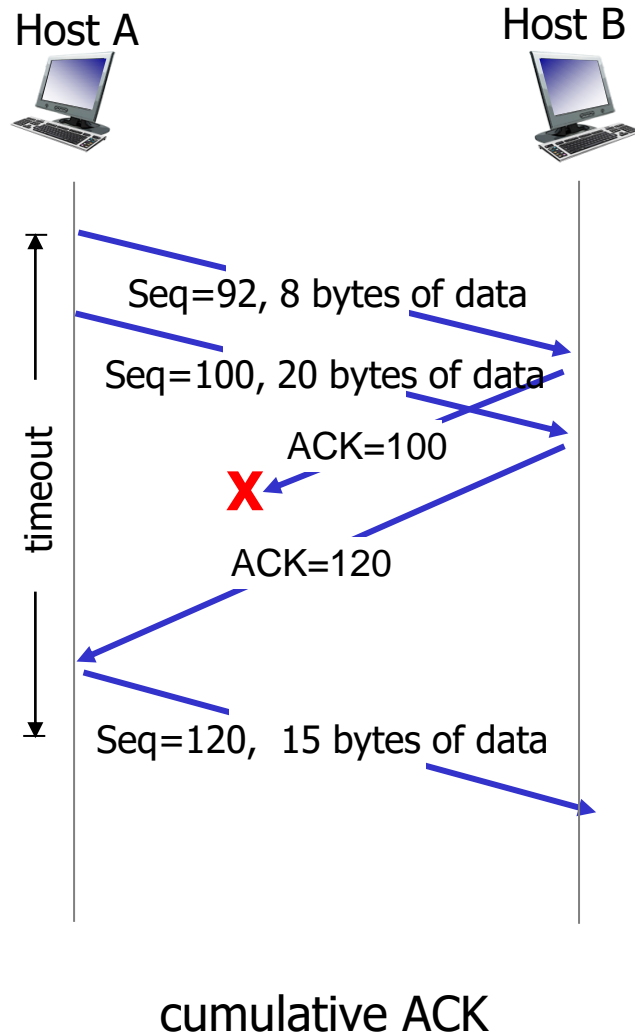


lost ACK scenario



premature timeout

# TCP: retransmission scenarios



# Doubling the Timeout Interval

length of the timeout interval after a timer expiration?

- ❖ whenever the timeout event occurs,
  - TCP retransmits the not-yet-acknowledged segment with the smallest sequence number
  - But each time TCP retransmits,
    - it sets the next timeout interval to twice the previous value

whenever the timer is started after

- either data received from application above,
- Or ACK received

Then TimeoutInterval is derived from the most recent values of EstimatedRTT and DevRTT

**The purpose is to control congestion**



# TCP fast retransmit


- ❖ time-out period often relatively long:
  - long delay before resending lost packet
- ❖ detect lost segments via duplicate ACKs.
  - sender often sends many segments back-to-back
  - if segment is lost, there will likely be many duplicate ACKs.

## *TCP fast retransmit*

if sender receives 3 ACKs for same data (“triple duplicate ACKs”), resend unacked segment with smallest seq #

- likely that unacked segment lost, so don't wait for timeout

# TCP fast retransmit

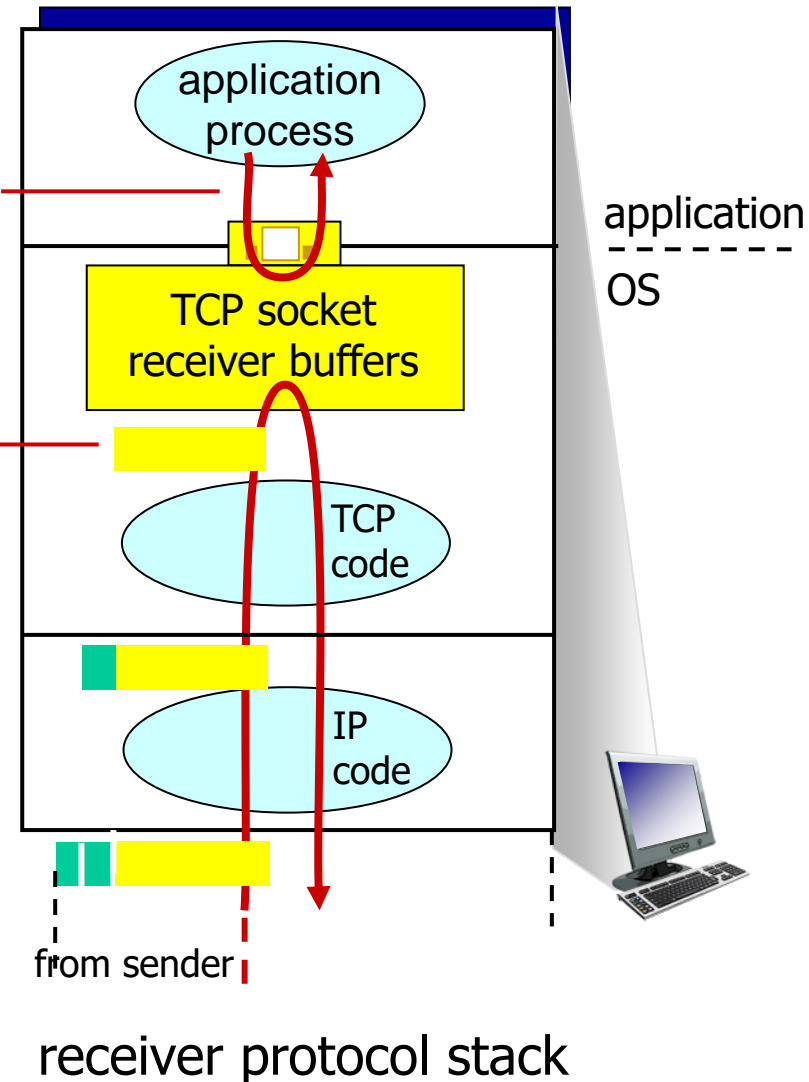


# TCP flow control

application may  
remove data from  
TCP socket buffers ....

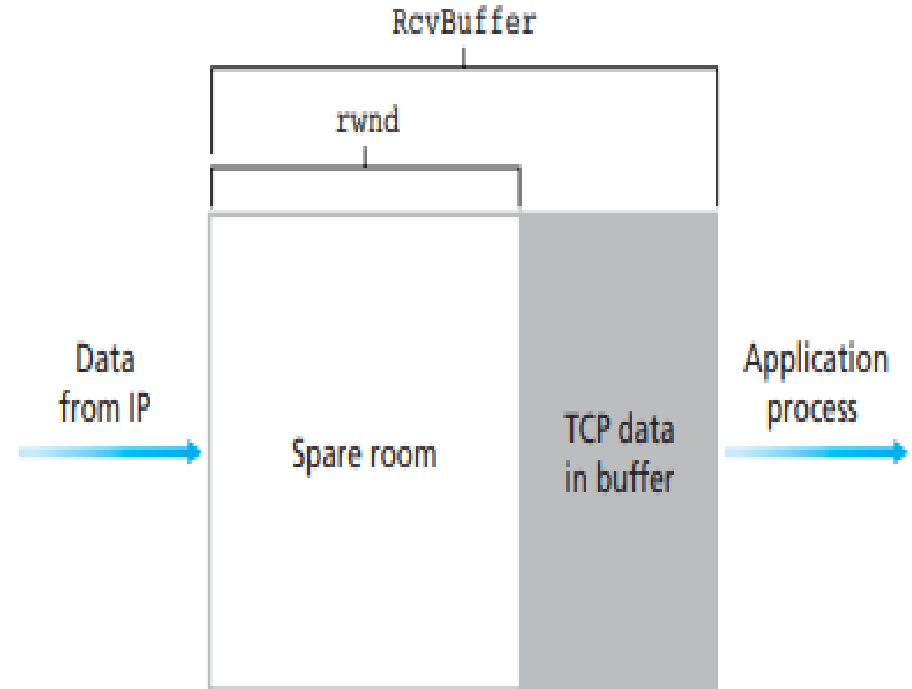
... slower than TCP  
receiver is delivering  
(sender is sending)

*flow control*  
receiver controls sender, so  
sender won't overflow  
receiver's buffer by transmitting  
too much, too fast



# TCP flow control

- ❖ receiver “advertises” free buffer space by including **rwnd** value in TCP header of receiver-to-sender segments
  - **RcvBuffer** size (typical default is 4096 bytes)
- ❖ sender limits amount of unacked (“in-flight”) data to receiver’s **rwnd** value
- ❖ guarantees receive buffer will not overflow



*receiver-side buffering*

# TCP flow control

**Receive Window is a variable maintained by receiver.**

$$\text{LastByteRcvd} - \text{LastByteRead} \leq \text{RcvBuffer}$$

if RcvBuffer is equal to left side equation, **buffer is full**  
else there is spare room

$$\text{rwnd} = \text{RcvBuffer} - [\text{LastByteRcvd} - \text{LastByteRead}]$$

$$\text{LastByteSent} - \text{LastByteAcked} \leq \text{rwnd}$$

**Question:** What happens when rwnd is 0!

# TCP flow control

**Question:** What happens when rwnd is 0!

- Application process at receiver empties the buffer and
- does not send new segment as receiver has nothing to send (neither data nor ACK).

So **sender** is unaware that space is there in receiver buffer and thus gets blocked

**Solution:**

- TCP specification requires sender to send segment with one data byte when rwnd is 0.
- Such segments will be ACKed by receiver and new ACKs will eventually contain non-zero rwnd values.

# Principles of congestion control

## *congestion:*

- ❖ informally: “too many sources sending too much data too fast for *network* to handle”
- ❖ different from flow control!
- ❖ manifestations:
  - lost packets (buffer overflow at routers)
  - long delays (queueing in router buffers)
- ❖ a top-10 problem!

# TCP congestion control

- TCP must use end-to-end congestion control rather than network-assisted congestion control
  - since the IP layer provides no explicit feedback to the end systems regarding network congestion
- TCP approach is to have each sender **limit the rate** at which it **sends traffic** into its **connection** as a function of **perceived** network **congestion**.



# TCP congestion control

- If there is **little** perceived congestion on the path between itself and the destination,
  - then the TCP sender **increases** its **send rate**;
- if the sender perceives that there is **more** congestion along the path,
  - then the sender **reduces** its **send rate**.

# TCP congestion control

- How does a TCP sender limit the rate at which it sends traffic into its connection?
- How does a TCP sender perceive that there is congestion on the path between itself and the destination?
- What algorithm should the sender use to change its send rate as a function of perceived end-to-end congestion?

# TCP congestion control

- **How does a TCP sender limit the rate at which it sends traffic into its connection?**
- In addition to send buffer, receive buffer and other variables like LastByteRead, rwnd etc)
- Additional variable: **congestion window** (cwnd)
  - Imposes a **constraint** on the send rate
- Specifically, the amount of **un-ACKed** data at a sender may not exceed the minimum of cwnd and rwnd
- **$\text{LastByteSent} - \text{LastByteAcked} \leq \min\{\text{cwnd}, \text{rwnd}\}$**


# TCP congestion control

- Assume TCP receive buffer is large enough,  
**ignore rwnd constraint**
  - So, the amount of **un-ACKed** data is solely limited by **cwnd**
- Further, assume sender always has data to send,
  - all segments in the **cwnd** are sent
- This constraint limits the amount of **un-ACKed** data at the sender
  - Thus indirectly limits the sender's send rate.

# TCP congestion control

- Consider a connection with negligible loss and packet transmission delays.
- At start of each RTT
  - the constraint permits the sender to **send cwnd bytes** of data into the connection;
- At the end of the RTT
  - the sender **receives ACKs** for the data.
- By adjusting the value of cwnd,
  - the sender can adjust send rate

# TCP Congestion Control



*TCP sending rate:*

- ❖ roughly: send `cwnd` bytes, wait RTT for ACKS, then send more bytes

$$\text{rate} \approx \frac{\text{cwnd}}{\text{RTT}} \text{ bytes/sec}$$

- ❖ sender limits transmission:

$$\text{LastByteSent} - \text{LastByteAcked} \leq \min\{\text{cwnd}, \text{rwnd}\}$$

- ❖ **cwnd** is dynamic, function of perceived network congestion

# TCP congestion control

- **How does a TCP sender perceives that there is congestion on the path?**
- **Loss event** at a sender
  - the occurrence of either a timeout or the receipt of three duplicate ACKs from the receiver.
- When there is excessive congestion, then one (or more) router buffers along the path overflows, causing a **datagram** (containing a TCP segment) to be **dropped**.
  - The dropped datagram, in turn, results in a **loss event** at the sender
    - either a timeout or the receipt of three duplicate ACKs
  - which is taken by the sender to be an indication of congestion on the sender-to-receiver path.

# TCP congestion control

- **Optimistic case:** Loss event does not occur
- ACKs for Un-ACKed data will be received at sender.
  - So ACKs indicate that all is well and increase its cwnd size (send rate)
- If ACKs arrive at a relatively slow rate,
  - then cwnd will be increased at a relatively slow rate.
- If ACKs arrive at a high rate, then the cwnd will be increased more quickly.
- Arrival of ACK is dependent on the end-end path and/or bandwidth of the link)
- TCP is self-clocking as it uses ACKs for increasing its cwnd size



# TCP congestion control

- **A lost segment implies congestion,**
  - rate should be decreased when a segment is lost.
- **An ACKed segment indicates all is well,**
  - rate can be increased when an original ACK arrives.
  -
- **Bandwidth probing:** how much increase or decrease?


# TCP congestion control

- **What algorithm should the sender use to change its send rate as a function of perceived end-to-end congestion?**

# TCP congestion control: additive increase multiplicative decrease

- ❖ *approach*: sender increases transmission rate (window size), probing for usable bandwidth, until loss occurs
  - *additive increase*: increase **cwnd** by 1 MSS every RTT until loss detected
  - *multiplicative decrease*: cut **cwnd** in half after loss

AIMD: probing  
for bandwidth



# TCP congestion control Algorithm


- Described [Jacobson 1988] and is standardized in [RFC 5681].

## **Three major components:**

1. slow start [Mandatory]
  2. congestion avoidance [Mandatory], and
  3. fast recovery [recommended but not required].
- Slow start and congestion avoidance differ in how they increase the size of cwnd in response to received ACKs.
    - slow start increases the size of cwnd more rapidly (despite its name) than congestion avoidance.

# TCP Slow Start

- ❖ when connection begins, increase rate exponentially until first loss event:
  - initially `cwnd` = 1 MSS
  - double `cwnd` every RTT
  - done by incrementing `cwnd` for every ACK received
- ❖ summary: initial rate is slow but ramps up exponentially fast



# TCP: detecting, reacting to loss

## ❖ **loss indicated by timeout:**

- **cwnd** set to 1 MSS;
- window then grows exponentially (as in slow start) to threshold, then grows linearly

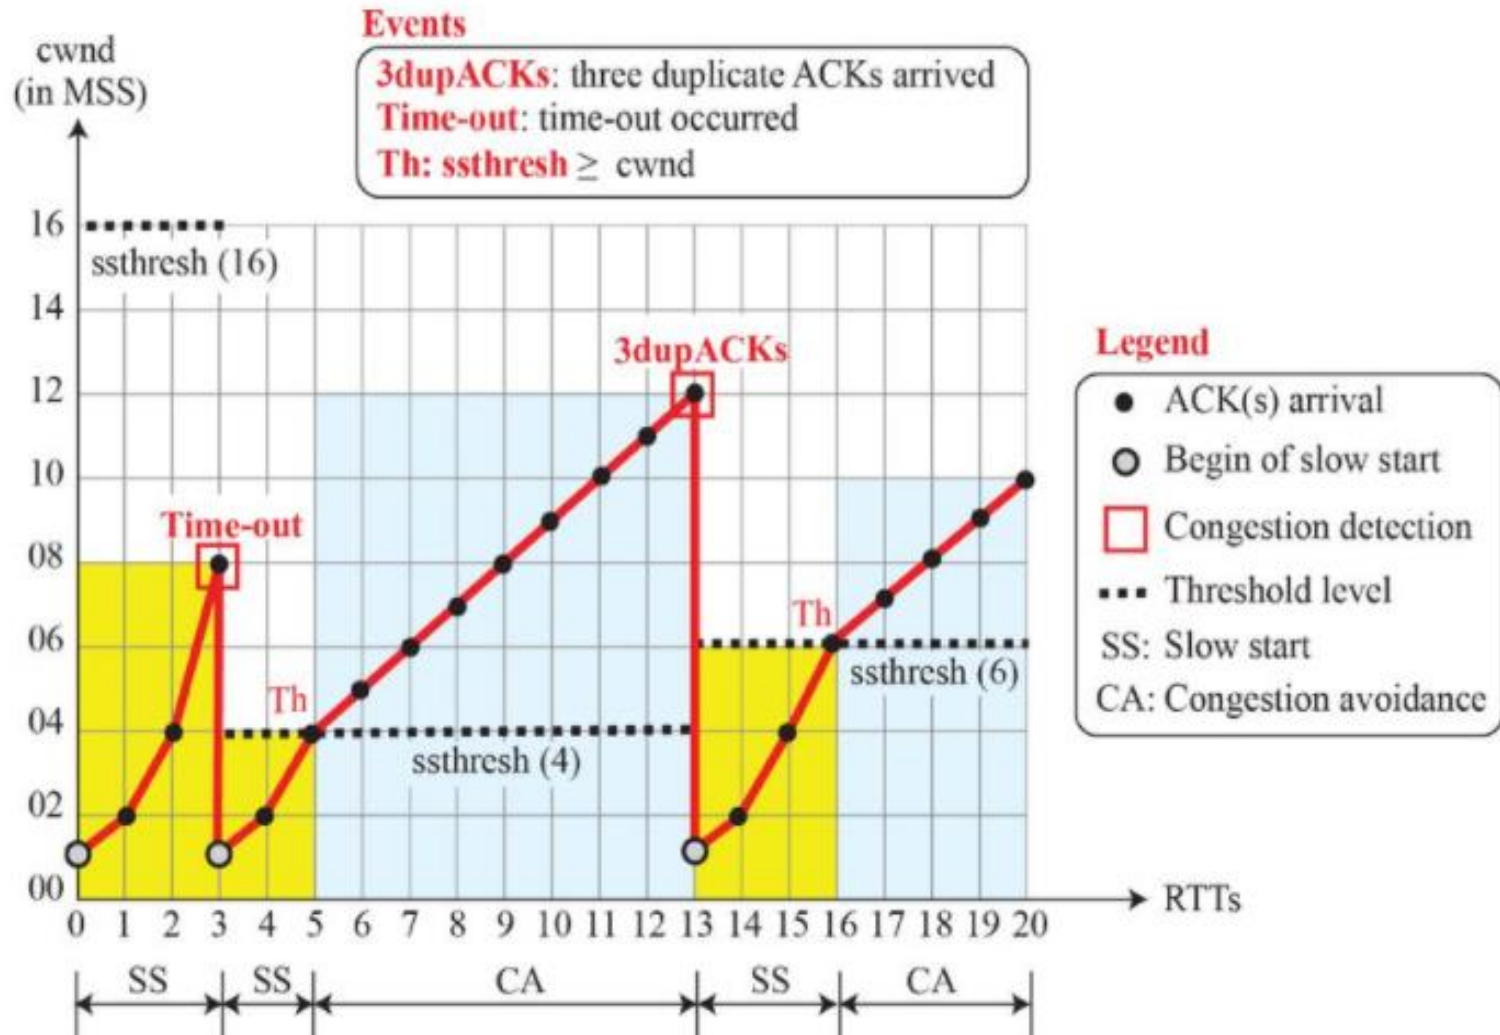
## ❖ **loss indicated by 3 duplicate ACKs:**

- ❖ TCP Tahoe (old TCP) always sets **cwnd** to 1 (timeout or 3 duplicate acks)

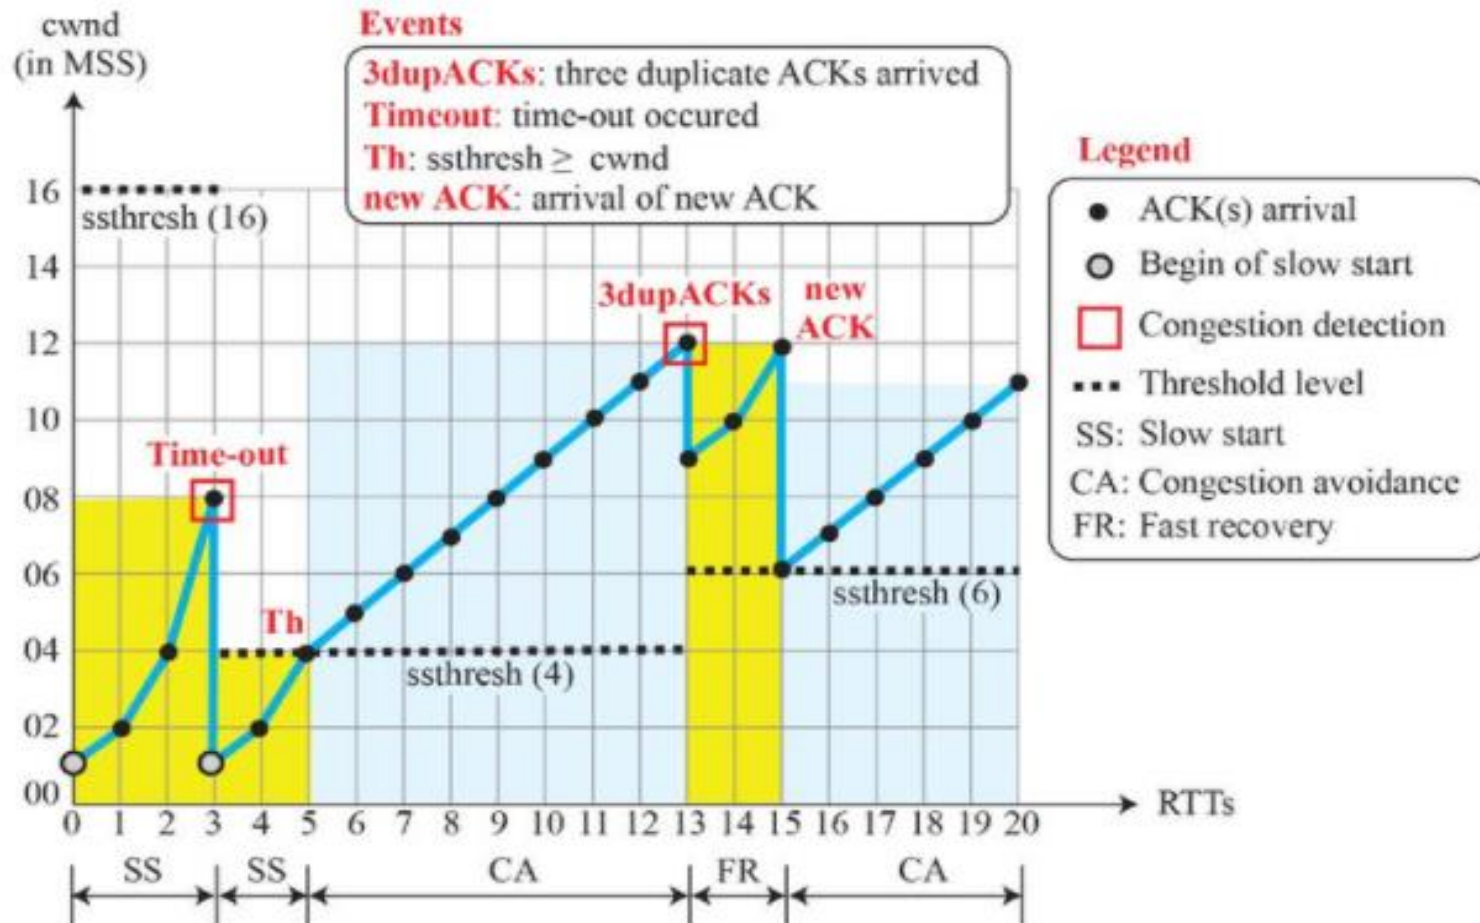
## ❖ TCP RENO

- dup ACKs indicate network capable of delivering some segments
- **cwnd** is cut in half window then grows linearly

# TCP Tahoe: Example



# TCP Reno: Example





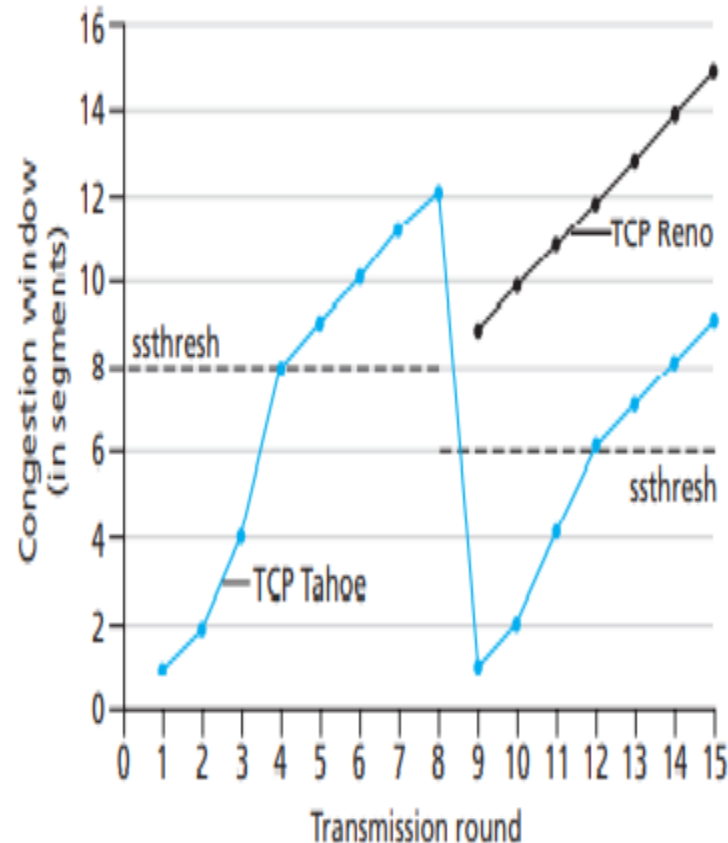
# TCP: switching from slow start to CA

**Q:** when should the exponential increase switch to linear?

**A:** when **cwnd** gets to 1/2 of its value before timeout.

## Implementation:

- ❖ variable **ssthresh**
- ❖ on loss event, **ssthresh** is set to 1/2 of **cwnd** just before loss event



# TCP: Reliable stream transport

- Reliability is achieved by buffering data combined with positive ACKs and retransmissions
- Before any data are transported, the two connection end points must explicitly set up a connection.
- A connection is identified by the IP addresses and TCP port numbers of the end points
- Many application layer protocols are defined over TCP,
  - such as HTTP (Web), SMTP (e-mail), and XMPP (instant messaging).